


# Recall

- Linear classifier  $h$
- 0-1 Loss
$$L(g, a) = \begin{cases} 0 & \text{if } g = a \\ 1 & \text{else} \end{cases}$$

- Training error
$$\mathcal{E}_n(h) = \frac{1}{n} \sum_{i=1}^n L(h(x^{(i)}), y^{(i)})$$



## A more-complete ML analysis

1. Establish a goal & find data
  - Example goal: diagnose whether people have heart disease based on their available information
2. Encode data in useful form for the ML algorithm
3. Run the ML algorithm & return a classifier
  - Example algorithms: (A) choose best classifier from a finite list; (B) perceptron; (C) averaged perceptron
4. Interpretation & evaluation

# A machine learning (ML) analysis

- First, need goal & data. E.g. diagnose whether people have heart disease based on their available information
- Next, put data in useful form for learning algorithm


	has heart disease?	resting heart rate (bpm)	pain?	job	medicines	age	family income (USD)
1	no	55	no	nurse	pain	40s	133000
2	no	71	no	admin	beta blockers, pain	20s	34000
3	yes	89	yes	nurse	beta blockers	50s	40000
4	no	67	no	doctor	none	50s	120000

# Encode data in usable form

- Identify the labels and encode as real numbers

has heart disease?	
1	no
2	no
3	yes
4	no

$$\{\text{'yes'}, \text{'no'}\} \leftrightarrow \{+1, -1\}$$



1	-1 $= y^{(1)}$
2	-1
3	+1
4	-1

- Depending on your algorithm, might instead use  $\{0, 1\}$
- Save mapping to recover predictions of new points


# Encode data in usable form

- Identify the features and encode as real numbers
- Feature: any function of the data (except labels)
- Today, old features:  $x$ ; new features:  $\phi(x)$

	resting heart rate (bpm)	pain?	job	medicines	age	family income (USD)
1	55	0	nurse	pain	40s	133000
2	71	0	admin	beta blockers, pain	20s	34000
3	89	1	nurse	beta blockers	50s	40000
4	67	0	doctor	none	50s	120000

# Encode categorical data


- Idea: turn each category into a unique natural number



# Encode categorical data

- Idea: turn each category into a unique binary number

	$\phi_d$	$\phi_{d+1}$	$\phi_{d+2}$
nurse	0	0	0
admin	0	0	1
pharmacist	0	1	0
doctor	0	1	1
social worker	1	0	0




# Encode categorical data

- Idea: turn each category into own unique 0-1 feature

	$\phi_d$	$\phi_{d+1}$	$\phi_{d+2}$	$\phi_{d+3}$	$\phi_{d+4}$
nurse	1	0	0	0	0
admin	0	1	0	0	0
pharmacist	0	0	1	0	0
doctor	0	0	0	1	0
social worker	0	0	0	0	1

- “one-hot encoding”



# Encode data in usable form

- Identify the features and encode as real numbers

	resting heart rate (bpm)	pain?	j1,j2,j3,j4,j5	medicines	age	family income (USD)
1	55	0	1,0,0,0,0	pain	40s	133000
2	71	0	0,1,0,0,0	beta blockers, pain	20s	34000
3	89	1	1,0,0,0,0	beta blockers	50s	40000
4	67	0	0,0,0,1,0	none	50s	120000




# Encode categorical data

- Should we use one-hot encoding?

	$\phi_d$	$\phi_{d+1}$	$\phi_{d+2}$	$\phi_{d+3}$
pain	1	0	0	0
pain & beta blockers	0	1	0	0
beta blockers	0	0	1	0
no medications	0	0	0	1


- Idea: factored encoding

	$\phi_d$	$\phi_{d+1}$
pain	1	0
pain & beta blockers	1	1
beta blockers	0	1
no medications	0	0



# Using a representative # for a range

- Potential pitfall: level of detail might be treated as meaningful (by you or others using the data)
- A way to diagnose many problems: plot your data!



# Encode data in usable form


- Identify the features and encode as real numbers

	resting heart rate (bpm)	pain?	j1,j2,j3,j4,j5	m1, m2	decade	family income (USD)
1	55	0	1,0,0,0,0	1,0	4	133000
2	71	0	0,1,0,0,0	1,1	2	34000
3	89	1	1,0,0,0,0	0,1	5	40000
4	67	0	0,0,0,1,0	0,0	5	120000

# Encode ordinal data

- Numerical data: order on data values, and differences in value are meaningful
- Categorical data: no order on data values
- Ordinal data: order on data values, but differences not meaningful
  - E.g. Likert scale:


Strongly disagree	Disagree	Neutral	Agree	Strongly agree
1	2	3	4	5



# Encode ordinal data

- Numerical data: order on data values, and differences in value are meaningful
- Categorical data: no order on data values
- Ordinal data: order on data values, but differences not meaningful
  - E.g. Likert scale:


Strongly disagree	Disagree	Neutral	Agree	Strongly agree
1,0,0,0,0	1,1,0,0,0	1,1,1,0,0	1,1,1,1,0	1,1,1,1,1



- Idea: Unary/ thermometer code


# Encode numerical data

- A closer look at the output of a linear classifier
- Idea: standardize numerical data
  - For  $d$ th feature:  $\phi_d^{(k)} = \frac{x_d^{(k)} - \text{mean}_d}{\text{stddev}_d}$




# Encode numerical data

- A closer look at the output of a linear classifier
- Idea: standardize numerical data
  - For  $d$ th feature:  $\phi_d^{(k)} = \frac{x_d^{(k)} - \text{mean}_d}{\text{stddev}_d}$



# More benefits of plotting your data

- And talking to experts









# Encode data in usable form

- Identify the features and encode as real numbers
- Standardize numerical features

	resting heart rate (bpm)	pain?	j1,j2,j3,j4,j5	m1, m2	decade	family income (USD)
1	-1.5	0	1,0,0,0,0	1,0	1	2.075
2	0.1	0	0,1,0,0,0	1,1	-1	-0.4
3	1.9	1	1,0,0,0,0	0,1	2	-0.25
4	-0.3	0	0,0,0,1,0	0,0	2	1.75

# Classification boundaries






# Nonlinear boundaries


- Idea: can approximate a smooth function with a  $k$ th order Taylor polynomial (e.g. around 0)

order ( $k$ )	terms when $d=1$	terms for general $d$
0	$[1]$	$[1]$
1	$[1, x_1]$	$[1, x_1, \dots, x_d]$
2	$[1, x_1, x_1^2]$	$[1, x_1, \dots, x_d, x_1^2, x_1 x_2, \dots, x_{d-1} x_d, x_d^2]$
3	$[1, x_1, x_1^2, x_1^3]$	$[1, x_1, \dots, x_d, x_1^2, x_1 x_2, \dots, x_{d-1} x_d, x_d^2, x_1^3, x_1^2 x_2, x_1 x_2 x_3, \dots, x_d^3]$

# Nonlinear boundaries



# Nonlinear boundaries



- Training error is 0!
- But seems like our classifier is overfitting

- How can we detect overfitting?
- How can we avoid overfitting?

# Evaluation of a learning algorithm


- How good is our learning algorithm on data like ours?

$x^{(1)}$


train

test

$x^{(n)}$
- Idea: use full data for training and then report training error
- Idea: reserve some data for testing
  - More training data: closer to training on full data
  - More testing data: less noisy estimate of performance
  - Only one classifier might not be representative
  - Good idea to shuffle order of data



# Evaluation of a learning algorithm



Cross-validate ( $\mathcal{D}_n$ ,  $k$ )

Divide  $\mathcal{D}_n$  into  $k$  chunks  $\mathcal{D}_{n,1}, \dots, \mathcal{D}_{n,k}$  (of roughly equal size)

**for**  $i = 1$  to  $k$

train  $h_i$  on  $\mathcal{D}_n \setminus \mathcal{D}_{n,i}$  (i.e. except chunk  $i$ )

compute “test” error  $\mathcal{E}(h_i, \mathcal{D}_{n,i})$  of  $h_i$  on  $\mathcal{D}_{n,i}$

**Return**  $\frac{1}{k} \sum_{i=1}^k \mathcal{E}(h_i, \mathcal{D}_{n,i})$

- Again, good idea to shuffle order of data first