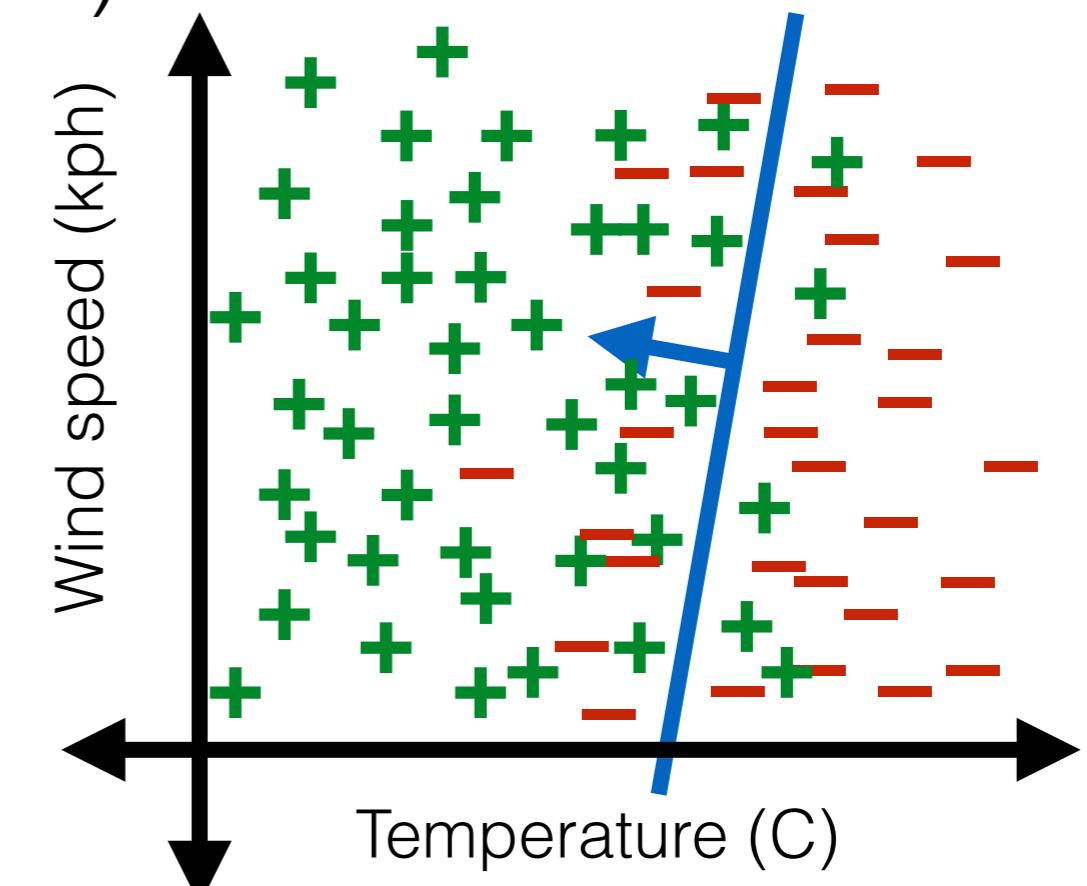
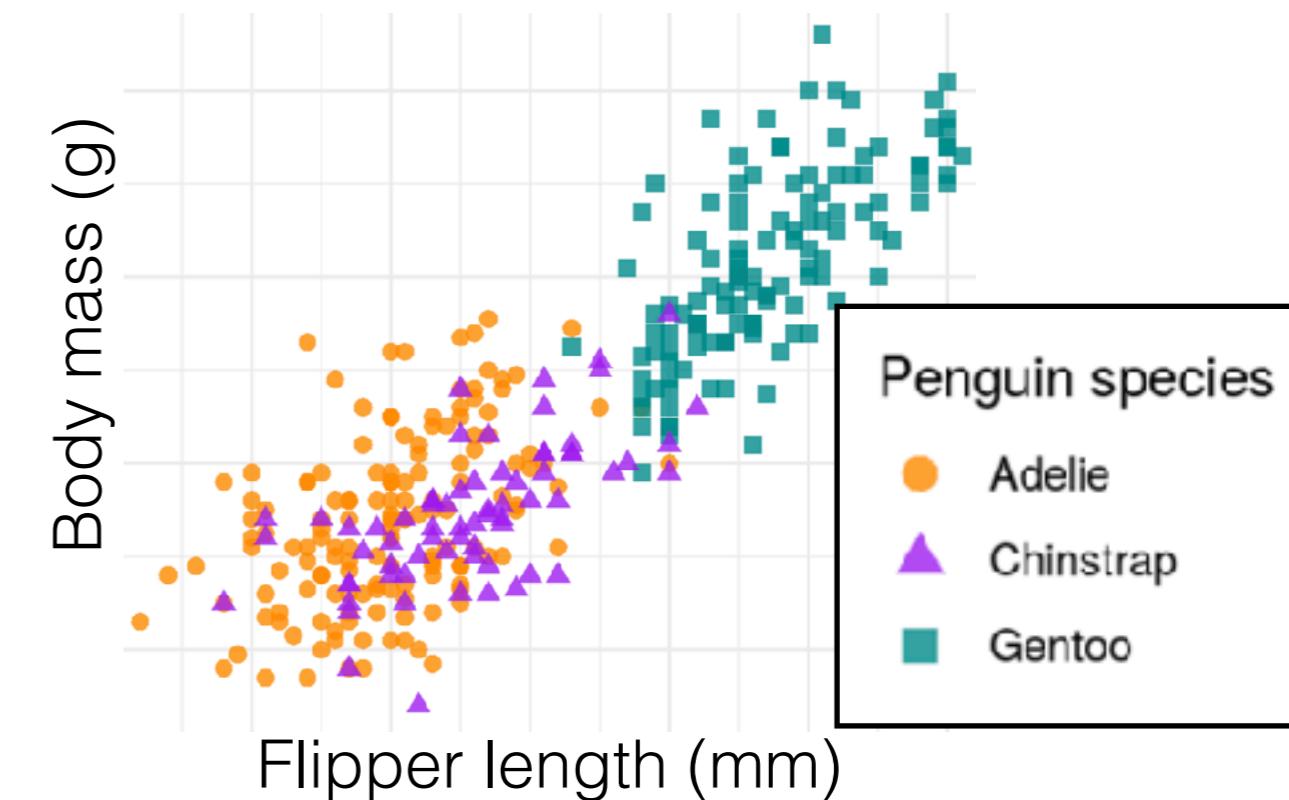
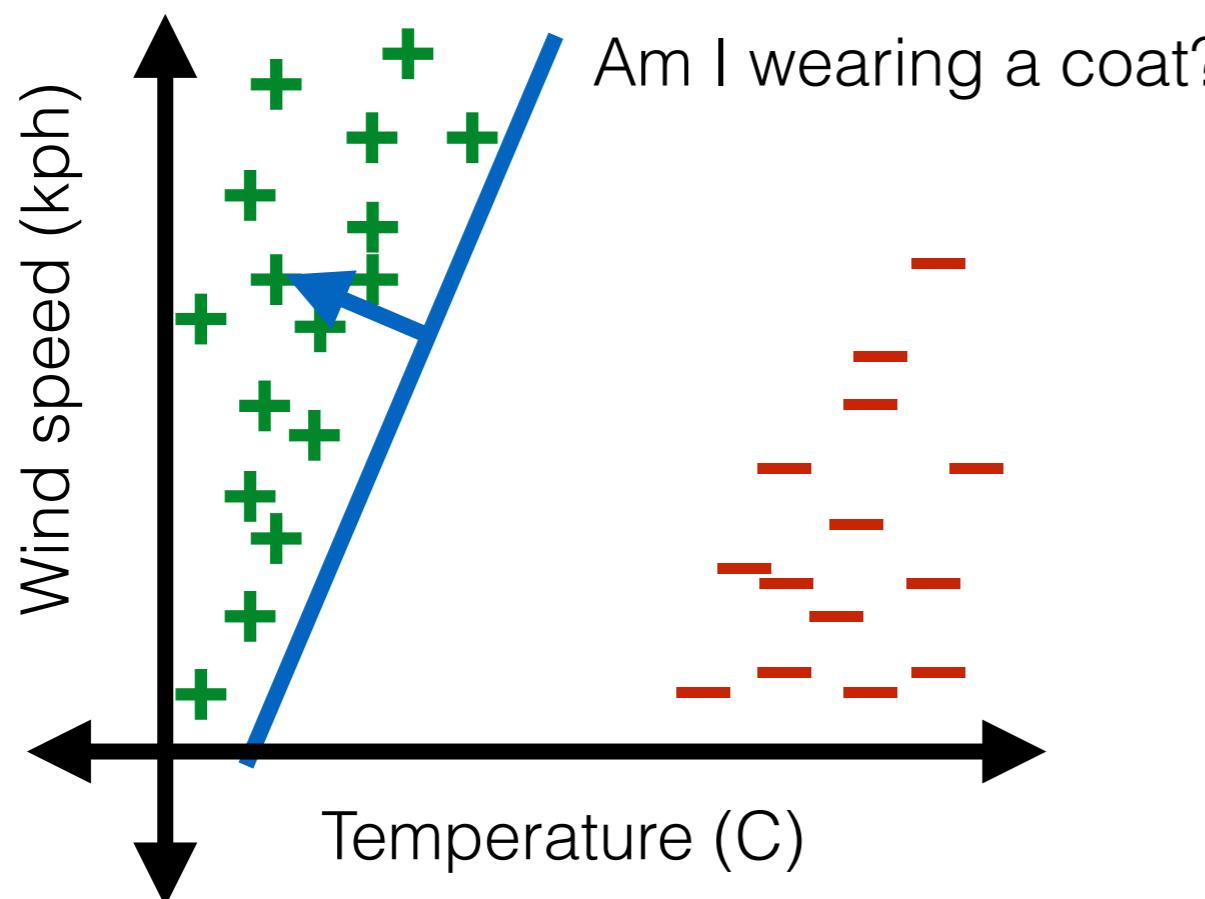


# Recall

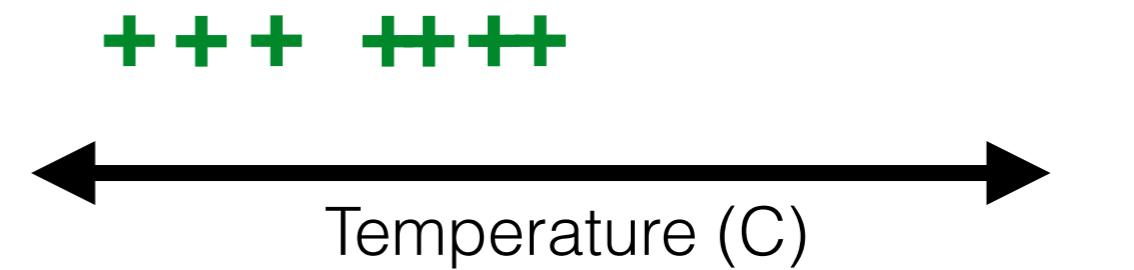
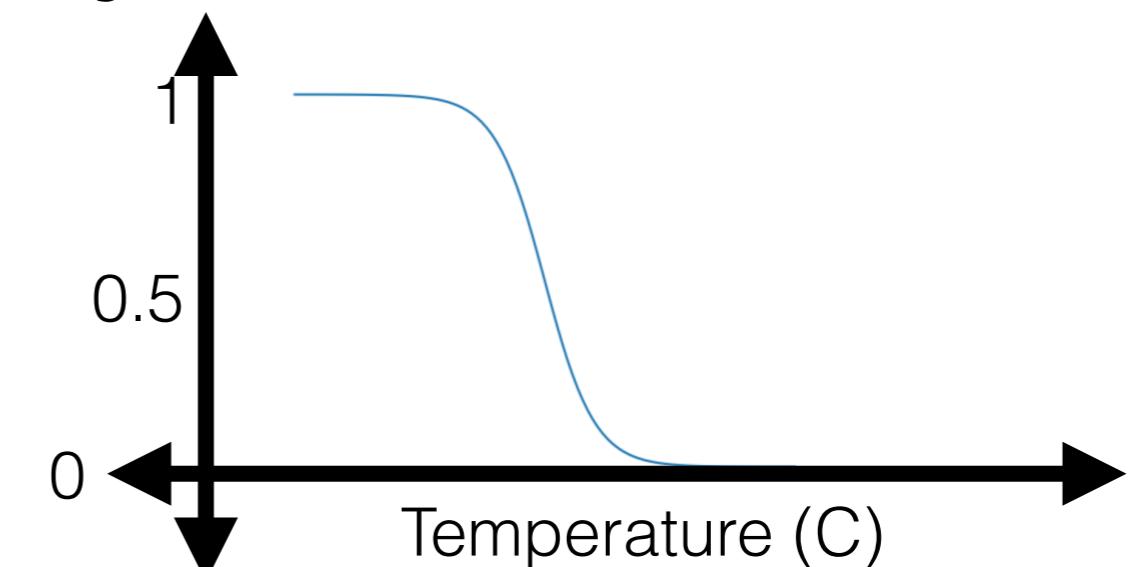
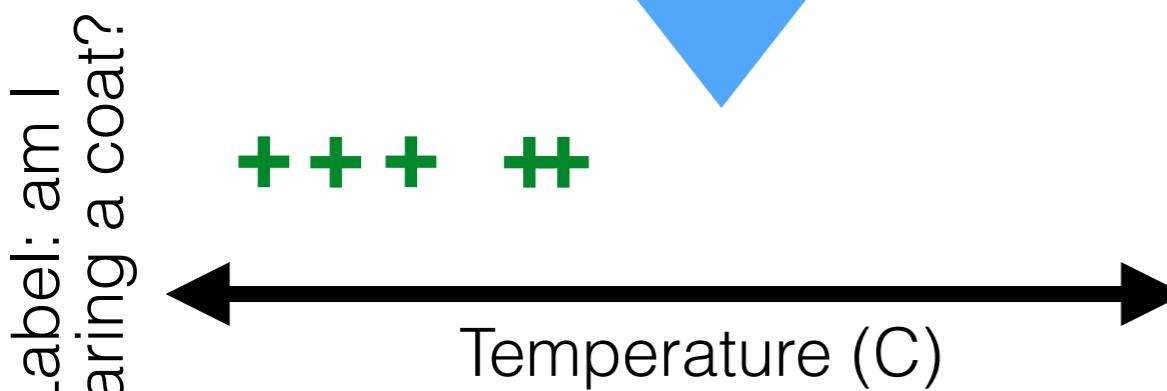
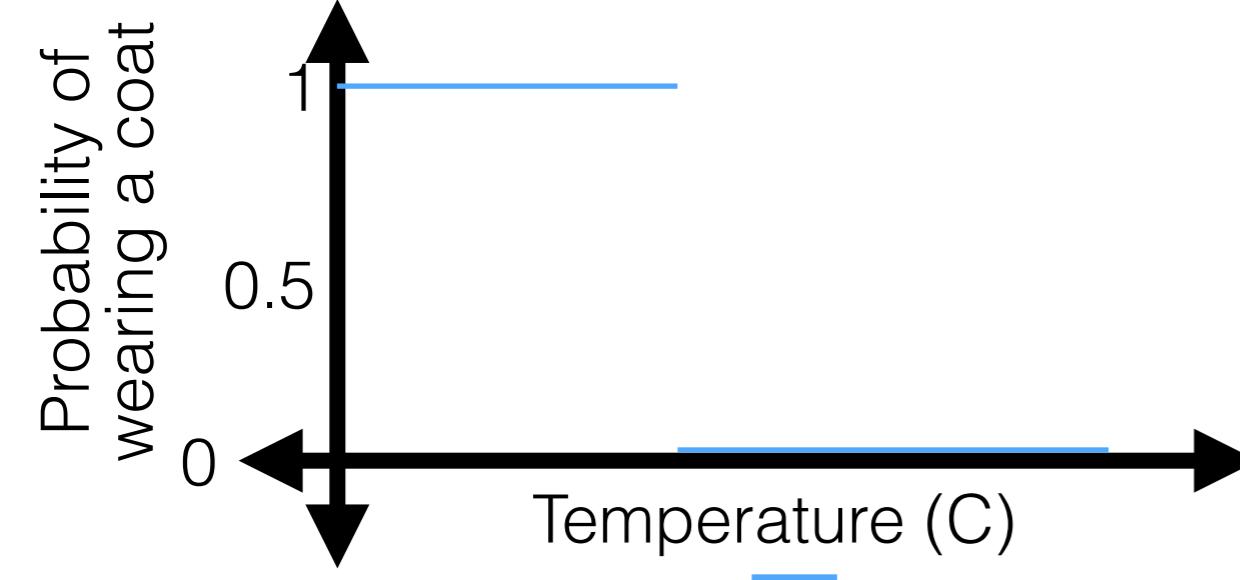
- Perceptron struggles with data that's not linearly separable

# Notice

- Perceptron doesn't have a notion of uncertainty (how well do we know what we know?)

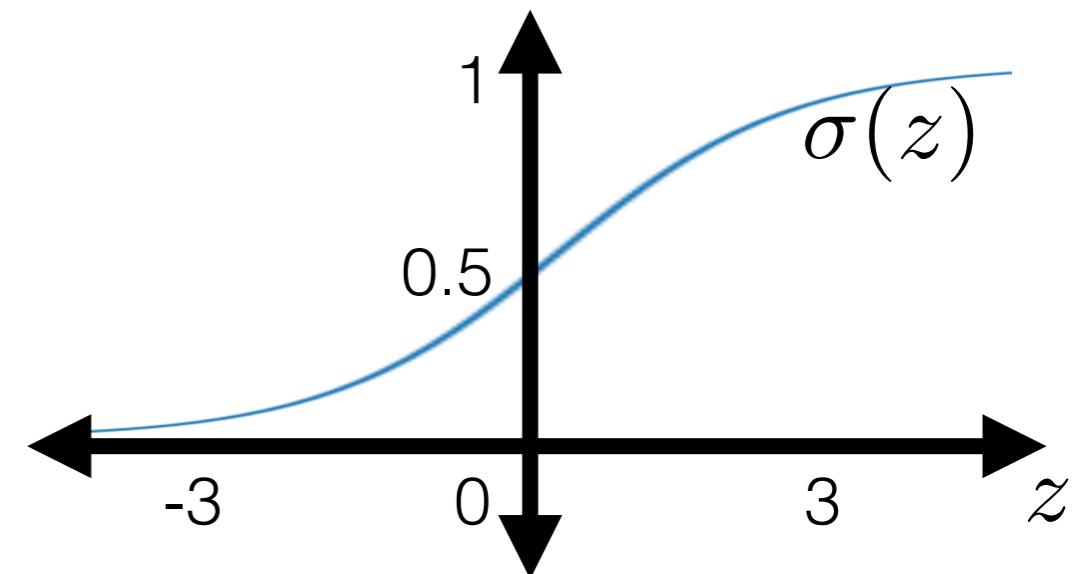


# Capturing uncertainty



- How to make this shape?
  - Sigmoid/logistic function

$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$



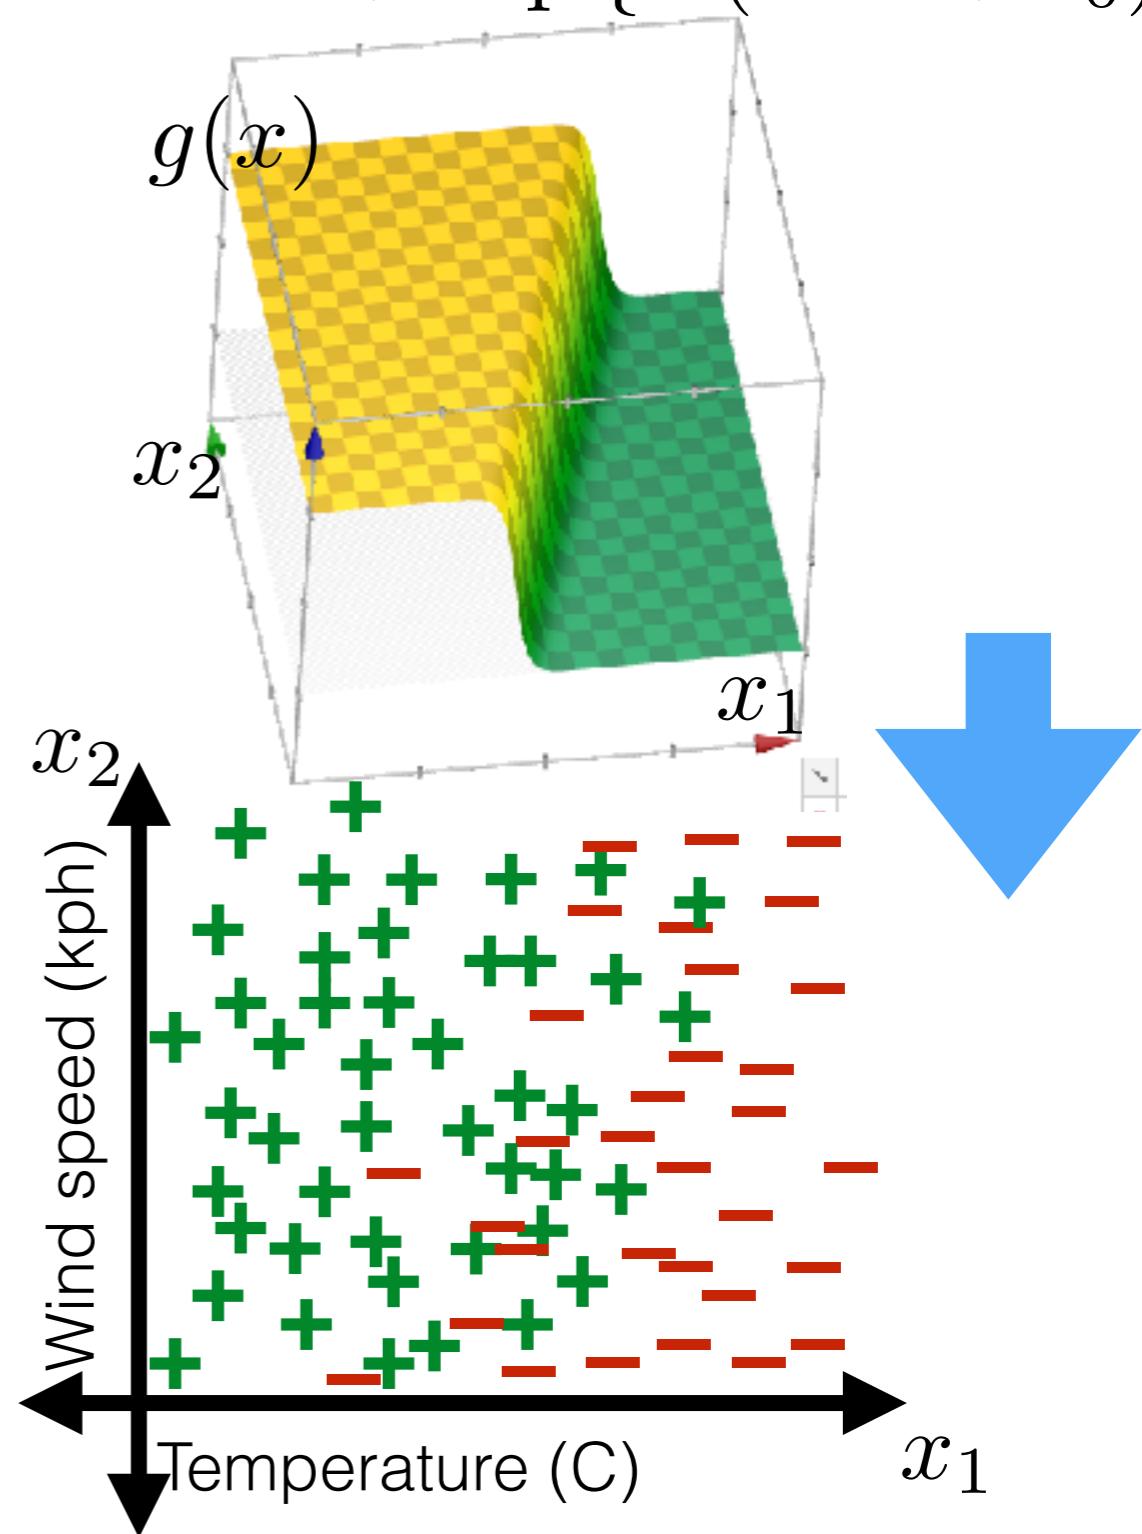
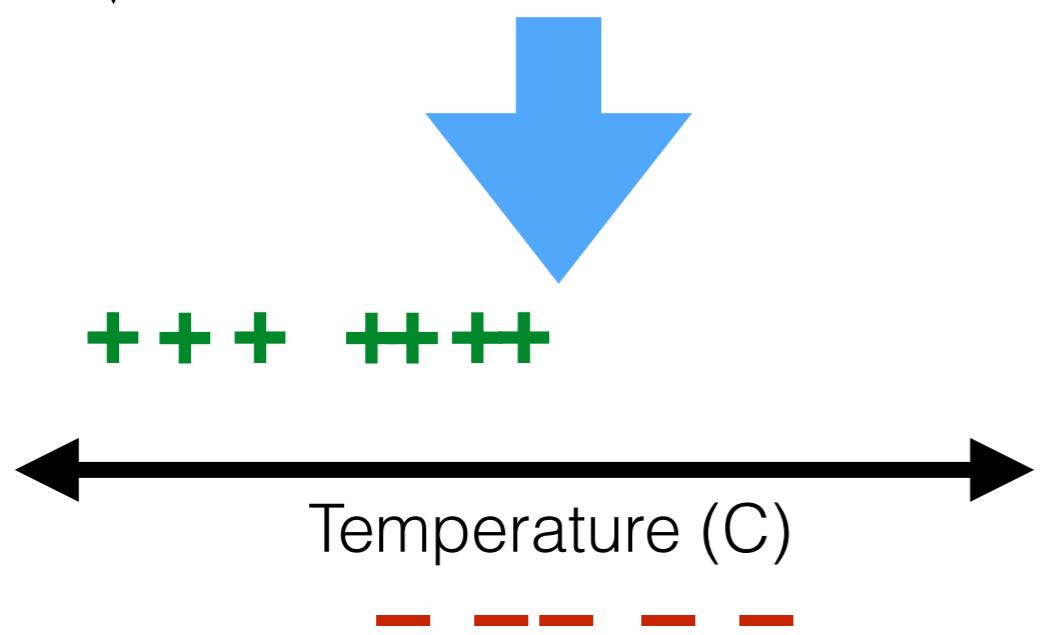
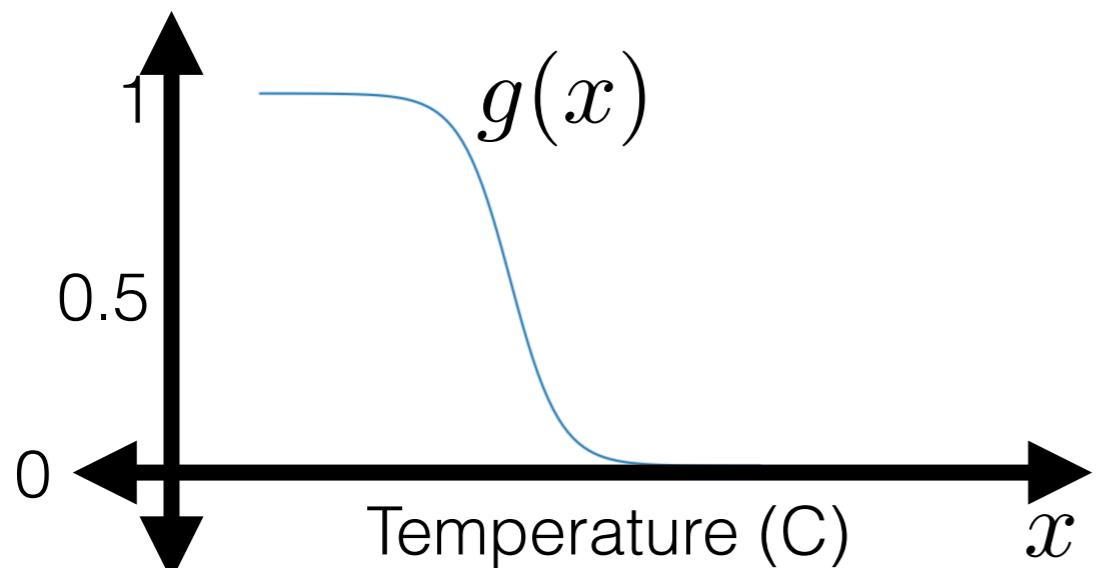
# Capturing uncertainty

2 features:

$$g(x) = \sigma(\theta^T x + \theta_0)$$
$$= \frac{1}{1 + \exp\{-(\theta^T x + \theta_0)\}}$$

1 feature:

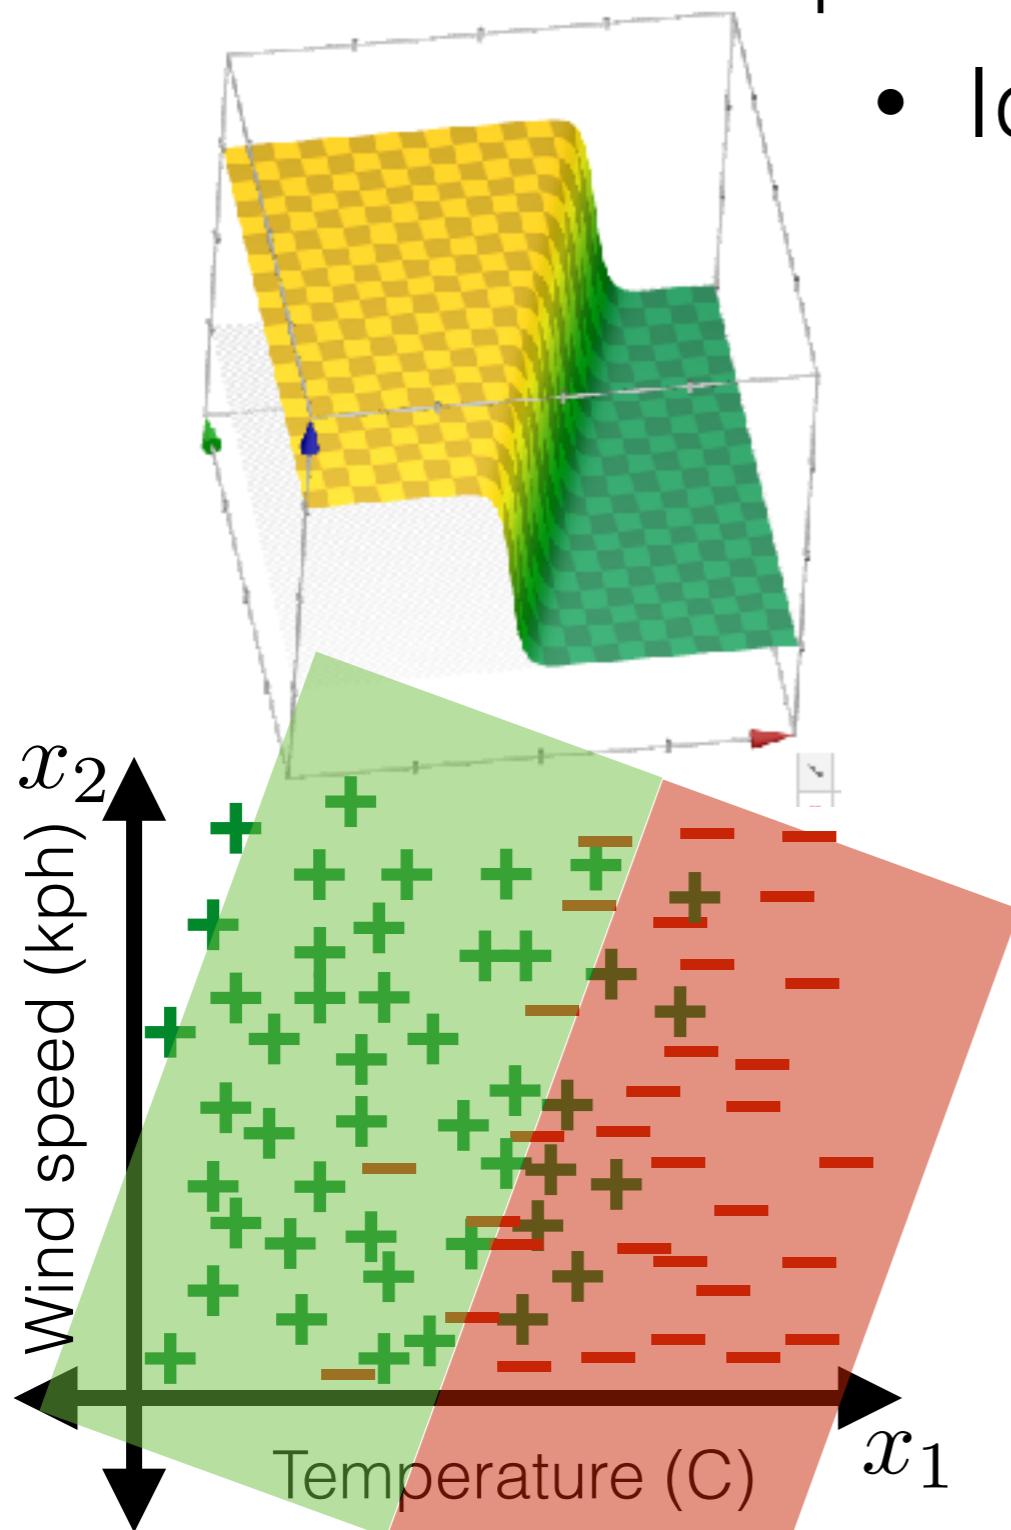
$$g(x) = \sigma(\theta x + \theta_0)$$
$$= \frac{1}{1 + \exp\{-(\theta x + \theta_0)\}}$$



# Linear logistic classification

aka logistic regression

- How do we learn a classifier (i.e. learn  $\theta, \theta_0$ )?
- How do we make predictions?



- Idea: predict +1 if: probability  $> 0.5$   
$$\sigma(\theta^\top x + \theta_0) > 0.5$$
  
$$\frac{1}{1 + \exp\{-(\theta^\top x + \theta_0)\}} > 0.5$$
  
$$\exp\{-(\theta^\top x + \theta_0)\} < 1$$
  
$$\theta^\top x + \theta_0 > 0$$

- Same hypothesis class as before! But we will get:
  - Uncertainties
  - Quality guarantees when data not linearly separable

# Linear logistic classification

aka logistic regression

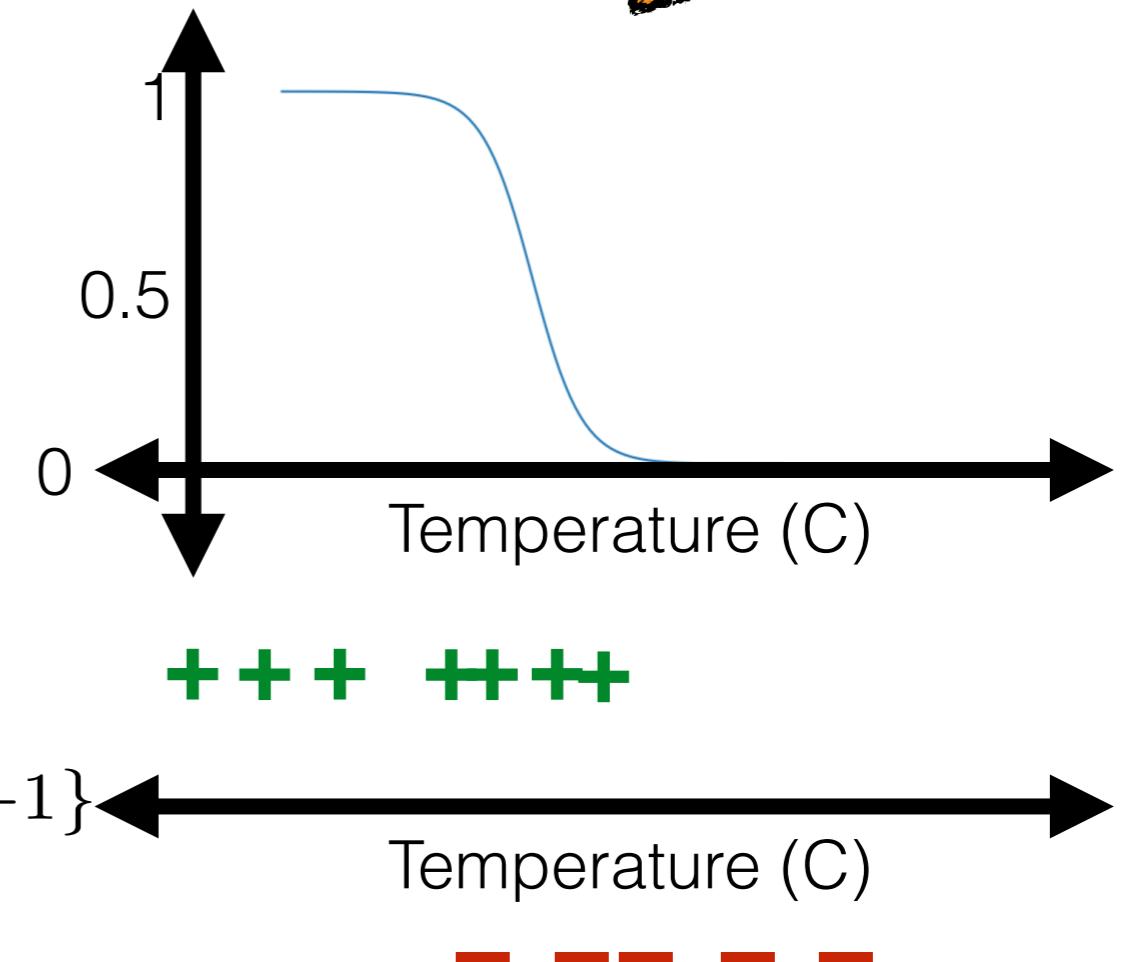
- How do we learn a classifier (i.e. learn  $\theta, \theta_0$ )?

Probability(data)

$$= \prod_{i=1}^n \text{Probability(data point } i) \\ \text{[Let } g^{(i)} = \sigma(\theta^\top x^{(i)} + \theta_0) \text{ ]}$$

$$= \prod_{i=1}^n \begin{cases} g^{(i)} & \text{if } y^{(i)} = +1 \\ (1 - g^{(i)}) & \text{else} \end{cases}$$

$$= \prod_{i=1}^n (g^{(i)})^{\mathbf{1}\{y^{(i)} = +1\}} (1 - g^{(i)})^{\mathbf{1}\{y^{(i)} \neq +1\}}$$



Loss(data) =  $-(1/n) * \log \text{probability(data)}$

$$= \frac{1}{n} \sum_{i=1}^n - \left( \mathbf{1}\{y^{(i)} = +1\} \log g^{(i)} + \mathbf{1}\{y^{(i)} \neq +1\} \log(1 - g^{(i)}) \right)$$

Negative log likelihood loss ( $g$  for guess,  $a$  for actual):

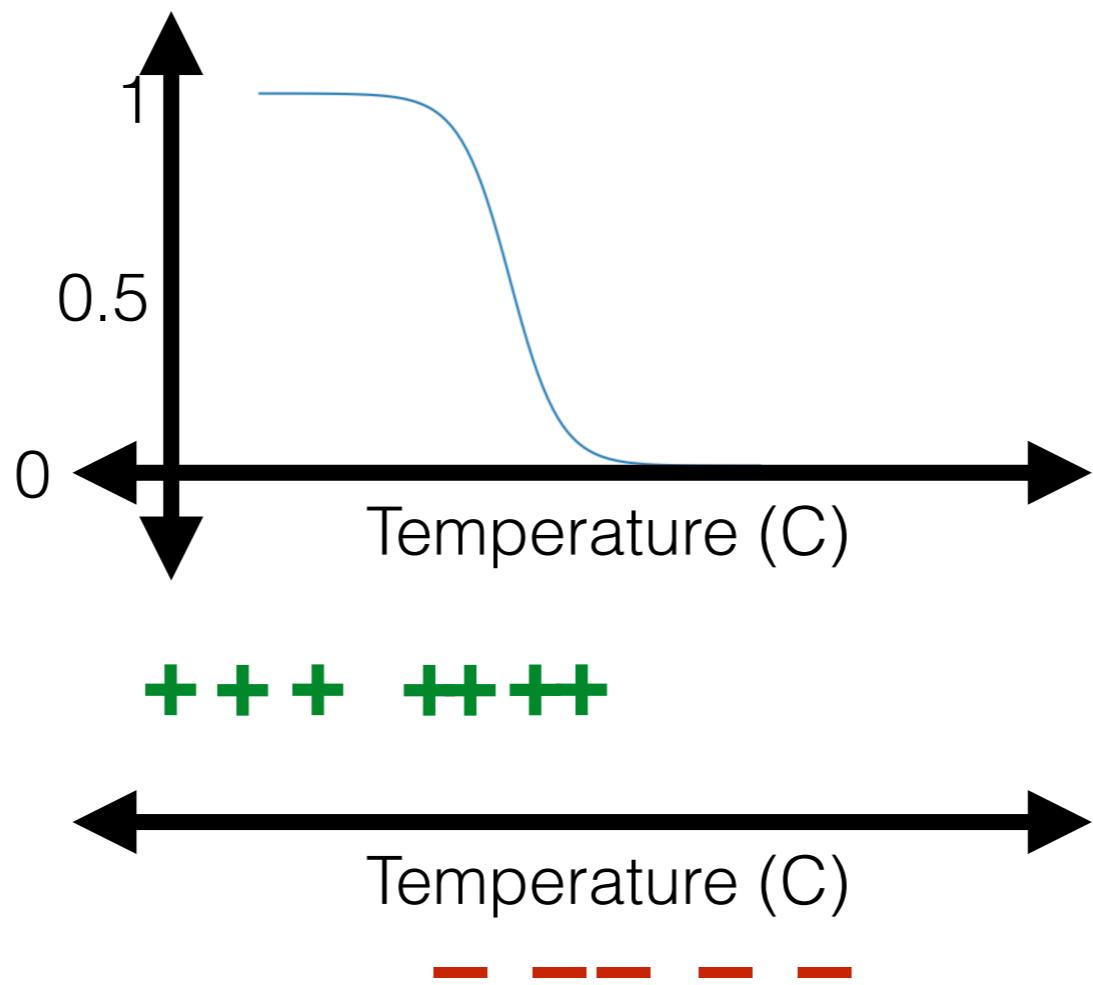
$$-L_{\text{nll}}(g, a) = (\mathbf{1}\{a = +1\} \log g + \mathbf{1}\{a \neq +1\} \log(1 - g))$$

# Linear logistic classification

aka logistic regression

- How do we learn a classifier (i.e. learn  $\theta, \theta_0$ )?
- Want to find parameter values to minimize average (negative log likelihood) loss across the data

$$J_{\text{lr}}(\Theta) = J_{\text{lr}}(\theta, \theta_0) = \frac{1}{n} \sum_{i=1}^n L_{\text{nll}}(\sigma(\theta^\top x^{(i)} + \theta_0), y^{(i)})$$



# Gradient descent

- Gradient  $\nabla_{\Theta} f = \left[ \frac{\partial f}{\partial \Theta_1}, \dots, \frac{\partial f}{\partial \Theta_m} \right]^T$ 
  - with  $\Theta \in \mathbb{R}^m$

Gradient-Descent ( $\Theta_{\text{init}}, \eta, f, \nabla_{\Theta} f, \epsilon$ )

Initialize  $\Theta^{(0)} = \Theta_{\text{init}}$

Initialize  $t = 0$

**repeat**

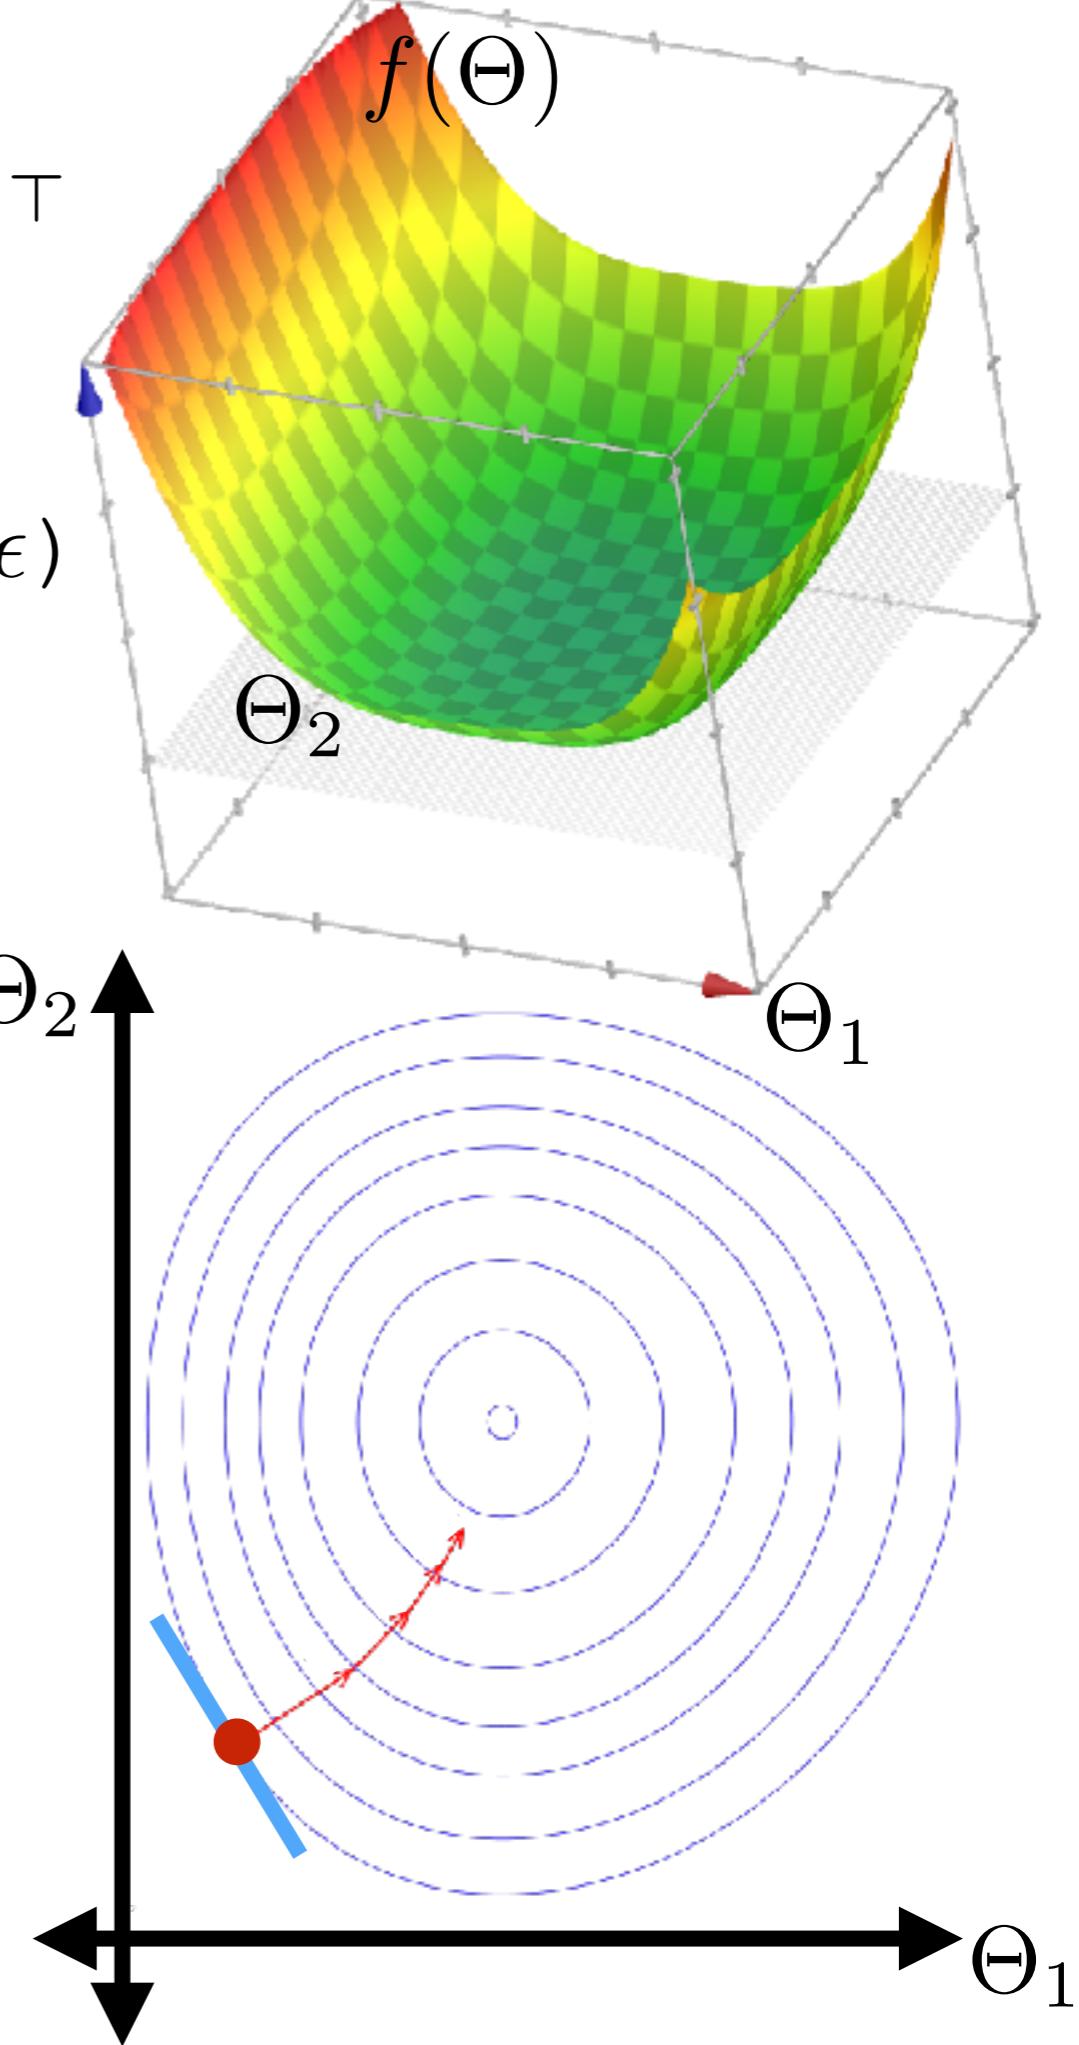
$t = t + 1$

$\Theta^{(t)} = \Theta^{(t-1)} - \eta \nabla_{\Theta} f(\Theta^{(t-1)})$

**until**  $|f(\Theta^{(t)}) - f(\Theta^{(t-1)})| < \epsilon$

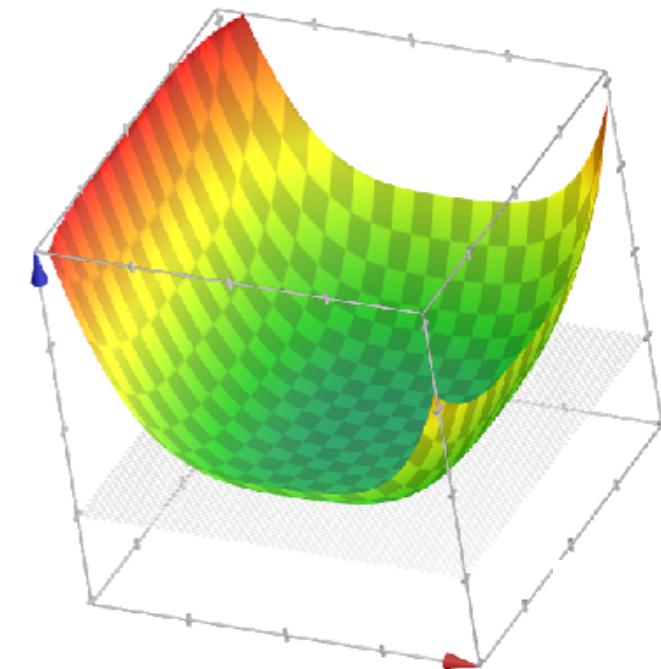
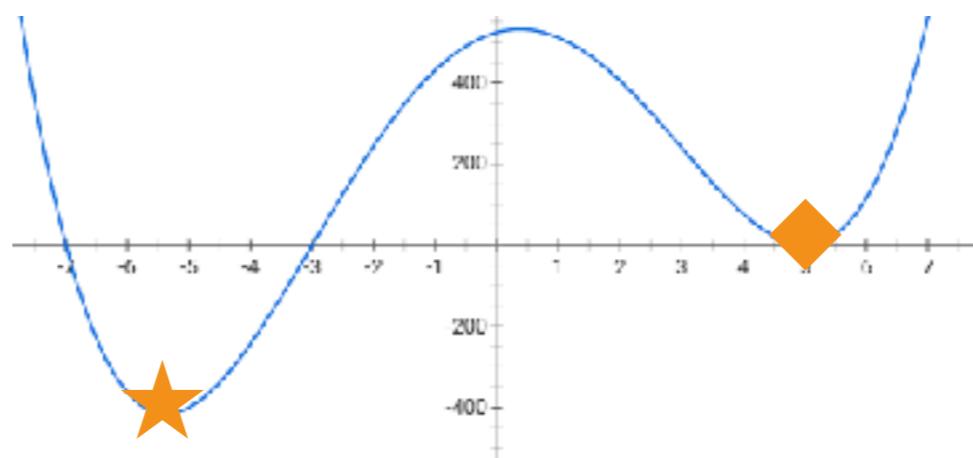
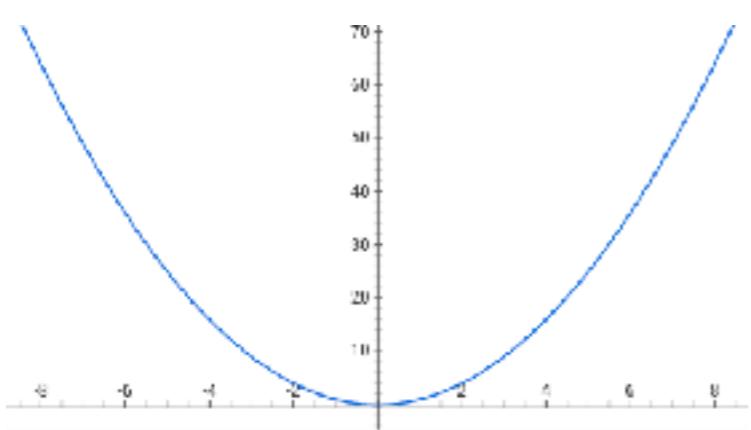
**Return**  $\Theta^{(t)}$

- Other possible stopping criteria:
  - Max number of iterations  $T$
  - $|\Theta^{(t)} - \Theta^{(t-1)}| < \epsilon$
  - $\|\nabla_{\Theta} f(\Theta^{(t)})\| < \epsilon$

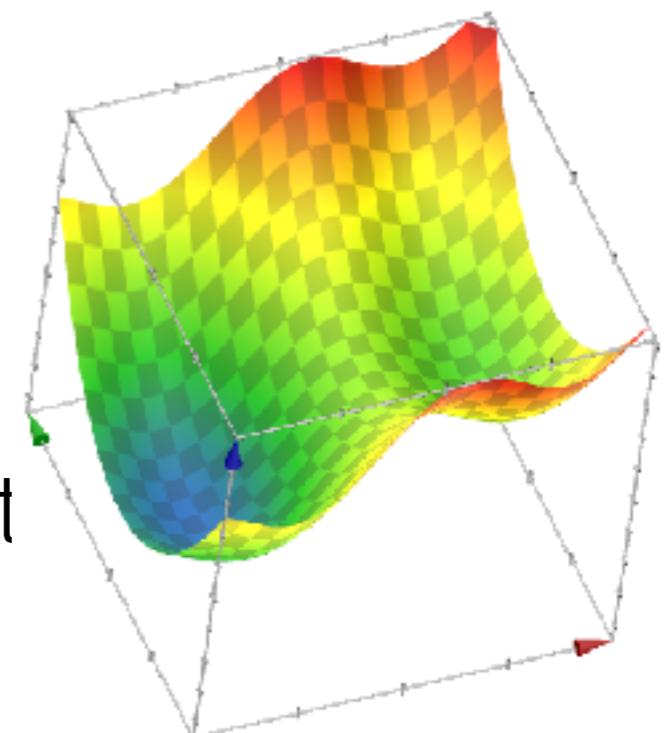


# Gradient descent properties

- A function  $f$  on  $\mathbb{R}^m$  is convex if any line segment connecting two points of the graph of  $f$  lies above or on the graph

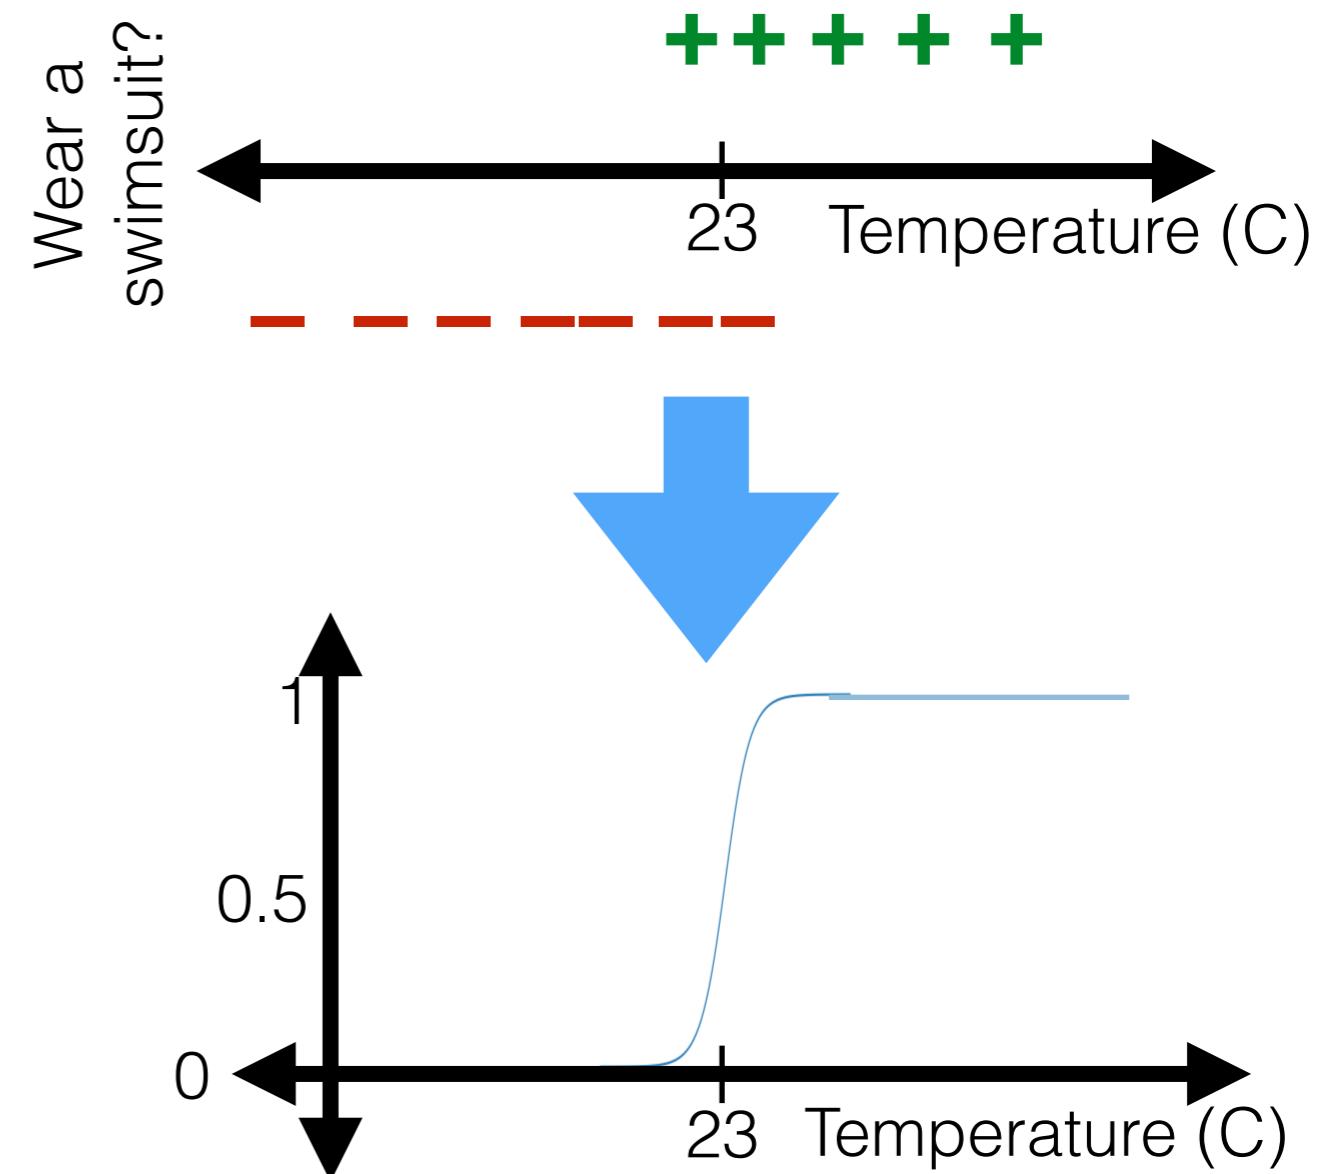
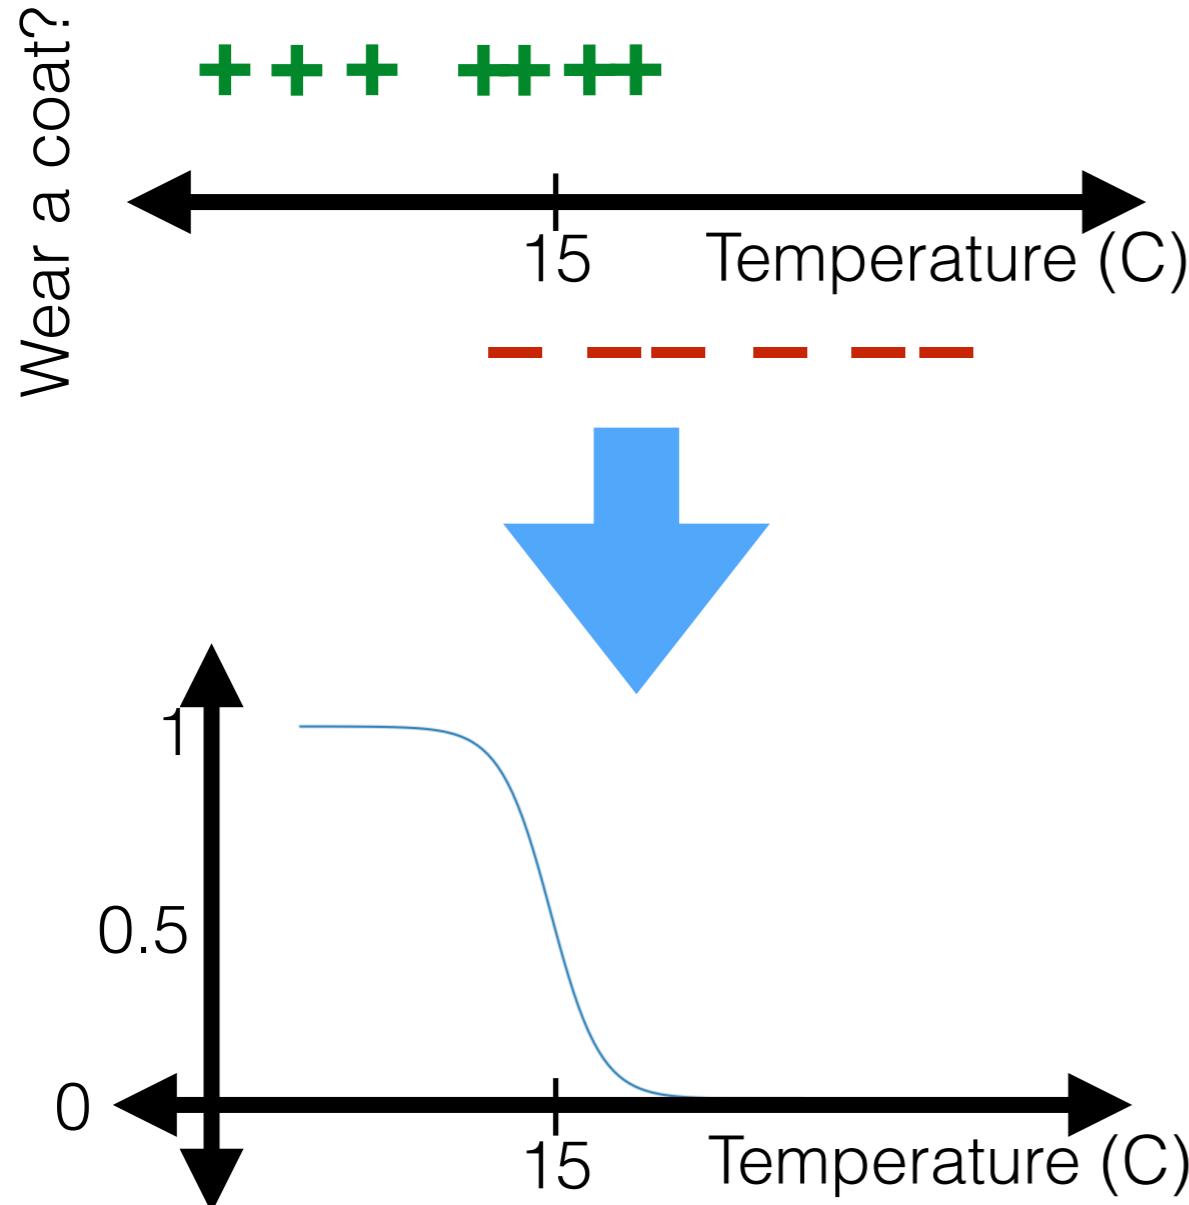


- **Theorem:** Gradient descent performance
  - **Assumptions:** (Choose any  $\tilde{\epsilon} > 0$ )
    - $f$  is sufficiently “smooth” and convex
    - $f$  has at least one global optimum
    - $\eta$  is sufficiently small
  - **Conclusion:** If run long enough, gradient descent will return a value within  $\tilde{\epsilon}$  of a global optimum  $\Theta$



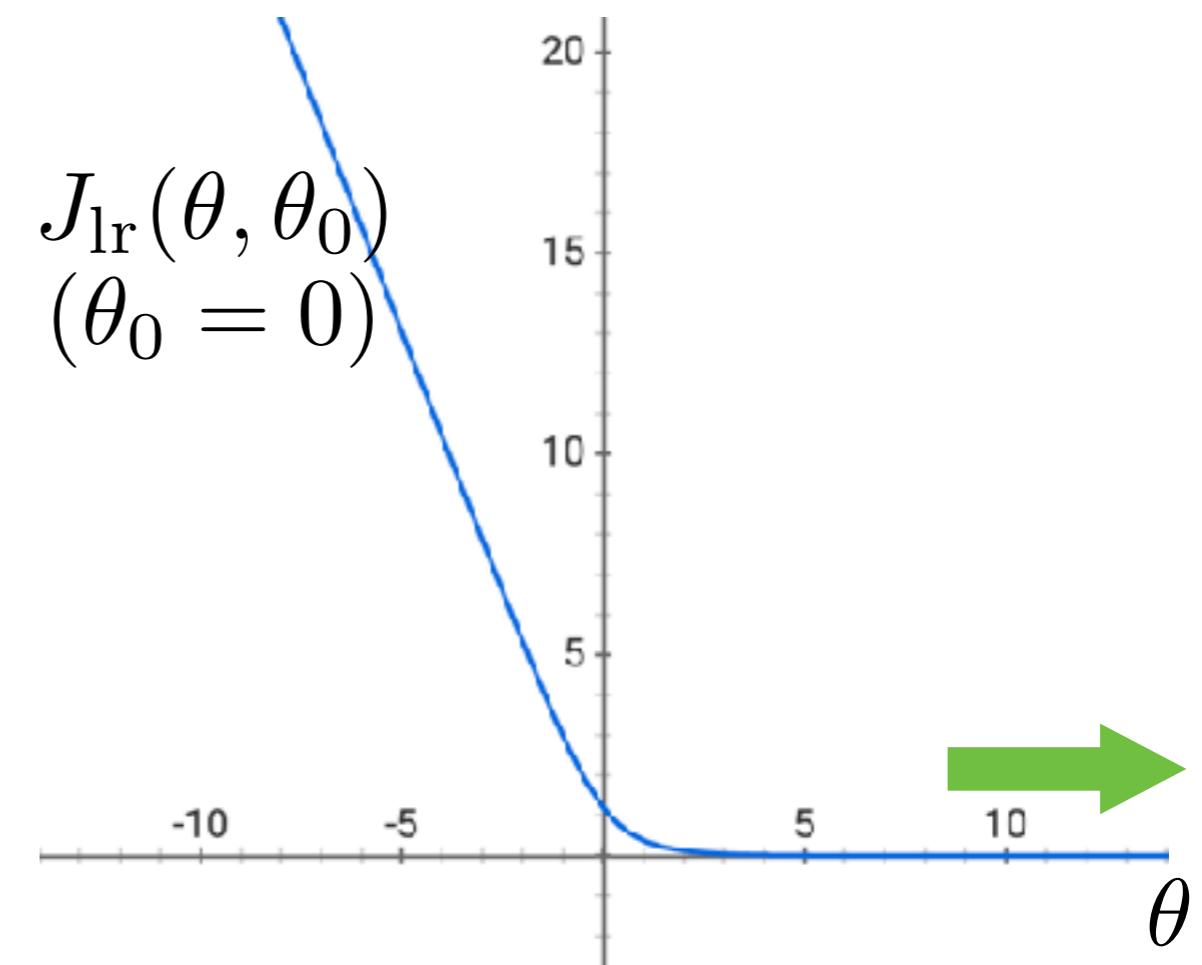
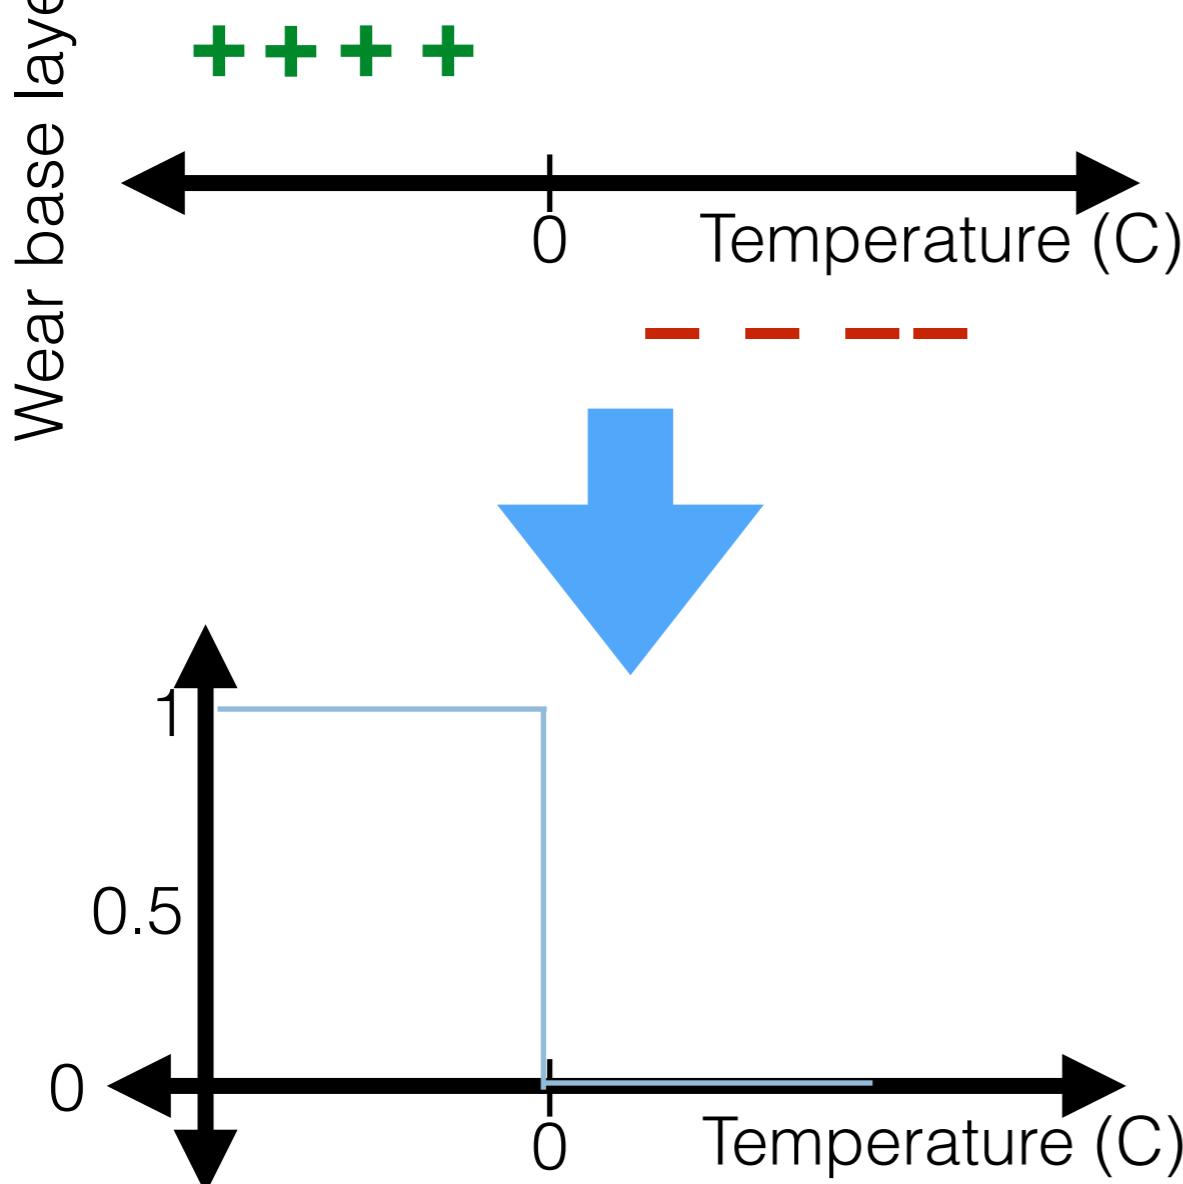
# Gradient descent for logistic regression

- Loss  $J_{\text{lr}}(\Theta) = J_{\text{lr}}(\theta, \theta_0)$  is differentiable & convex
- Run Gradient-Descent ( $\Theta_{\text{init}}, \eta, J_{\text{lr}}, \nabla_{\Theta} J_{\text{lr}}, \epsilon$ )



# Gradient descent for logistic regression

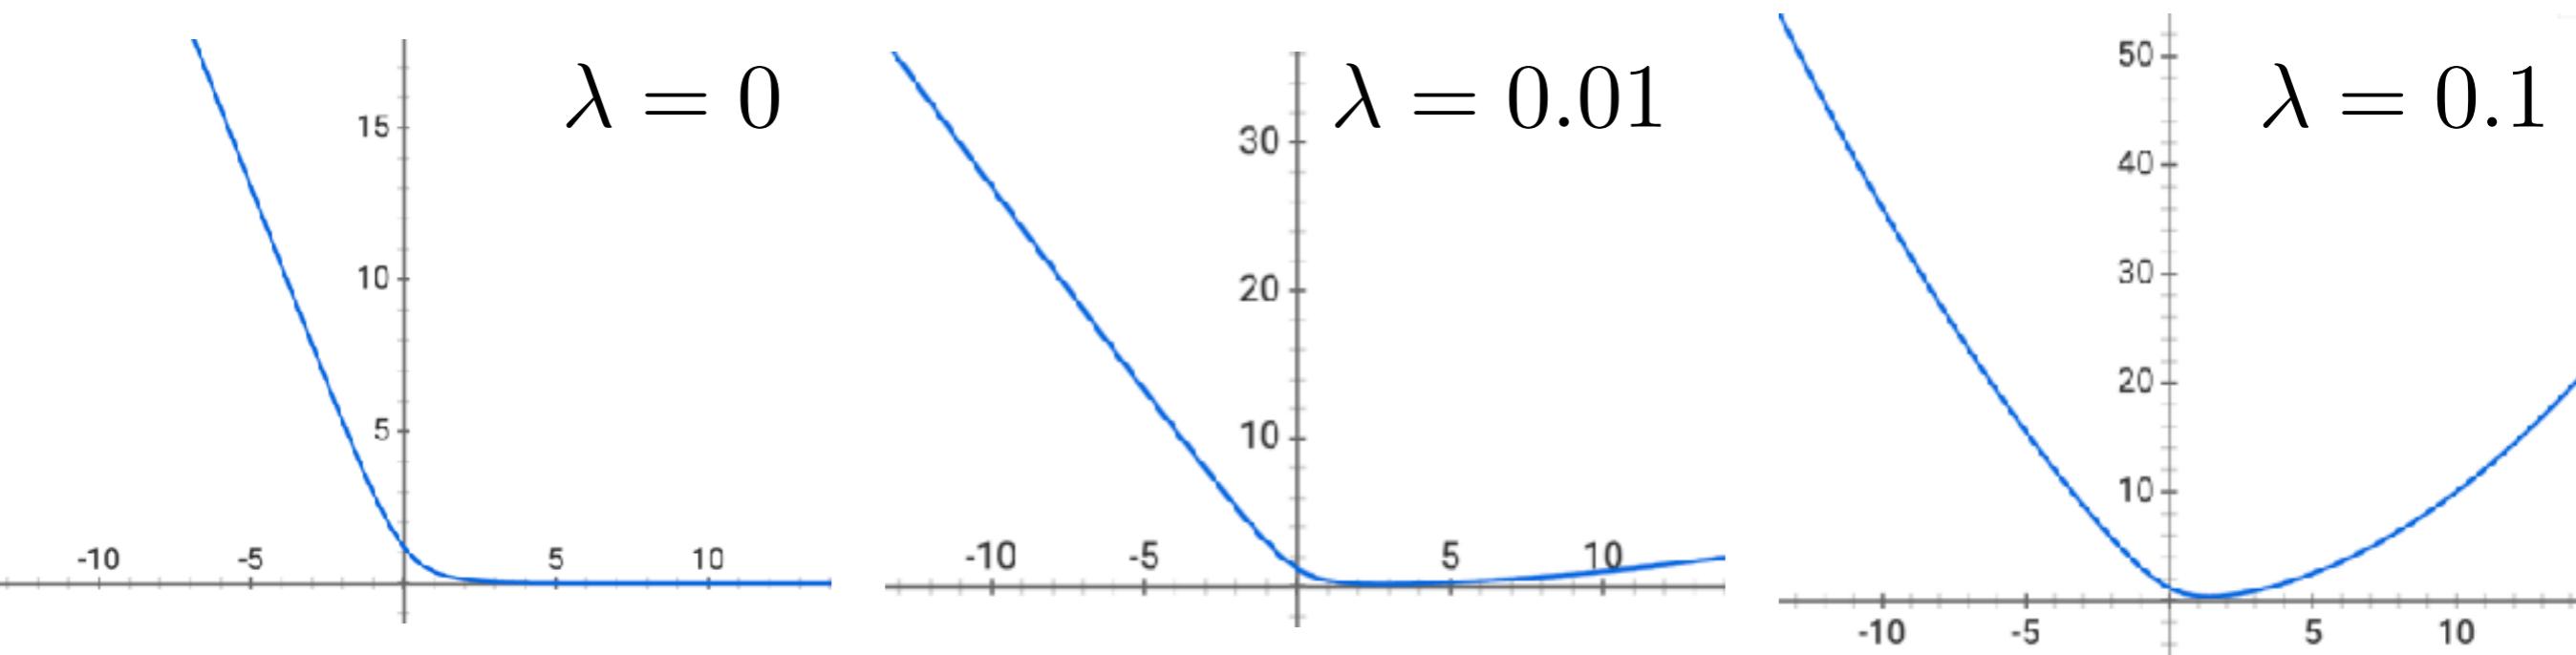
- Loss  $J_{\text{lr}}(\Theta) = J_{\text{lr}}(\theta, \theta_0)$  is differentiable & convex
- Run Gradient-Descent ( $\Theta_{\text{init}}, \eta, J_{\text{lr}}, \nabla_{\Theta} J_{\text{lr}}, \epsilon$ )



# Logistic regression loss revisited

$$\begin{aligned} J_{\text{lr}}(\Theta) &= J_{\text{lr}}(\theta, \theta_0) \\ &= \frac{1}{n} \sum_{i=1}^n L_{\text{nll}}(\sigma(\theta^\top x^{(i)} + \theta_0), y^{(i)}) + \lambda \|\theta\|^2 \quad (\lambda \geq 0) \end{aligned}$$

- A “regularizer” or “penalty”  $R(\theta) = \lambda \|\theta\|^2$
- Penalizes being overly certain
- Objective is still differentiable & convex (gradient descent)



- How to choose hyperparameters? One option: consider a handful of possible values and compare via CV

# Logistic regression learning algorithm

LR-Gradient-Descent ( $\theta_{\text{init}}, \theta_{0,\text{init}}, \eta, \epsilon$ )

Initialize  $\theta^{(0)} = \theta_{\text{init}}$

Initialize  $\theta_0^{(0)} = \theta_{0,\text{init}}$

Initialize  $t = 0$

**repeat**

$t = t + 1$

$$\theta^{(t)} = \theta^{(t-1)} - \eta \left\{ \frac{1}{n} \sum_{i=1}^n [\sigma(\theta^{(t-1)\top} x^{(i)} + \theta_0^{(t-1)}) - y^{(i)}] x^{(i)} + 2\lambda\theta^{(t-1)} \right\}$$

$$\theta_0^{(t)} = \theta_0^{(t-1)} - \eta \left\{ \frac{1}{n} \sum_{i=1}^n [\sigma(\theta^{(t-1)\top} x^{(i)} + \theta_0^{(t-1)}) - y^{(i)}] \right\}$$

**until**  $|J_{\text{lr}}(\theta^{(t)}, \theta_0^{(t)}) - J_{\text{lr}}(\theta^{(t-1)}, \theta_0^{(t-1)})| < \epsilon$

**Return**  $\theta^{(t)}, \theta_0^{(t)}$

Exactly gradient descent  
with  $f$  given by logistic  
regression objective