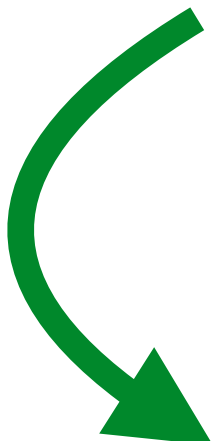





Images



1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

- We'll focus on grayscale images
 - Each pixel takes a value between 0 and P
 - Here, 0: black, 1: white
 - Larger P in Lab Week 08
- How do we use an image as an input for a neural net?


Images



- We'll focus on grayscale images
 - Each pixel takes a value between 0 and P
 - Here, 0: black, 1: white
 - Larger P in Lab Week 08
- How do we use an image as an input for a neural net?

Previous neural nets in this class

- Recall: *Fully connected layer: every input is connected to every output by a weight*



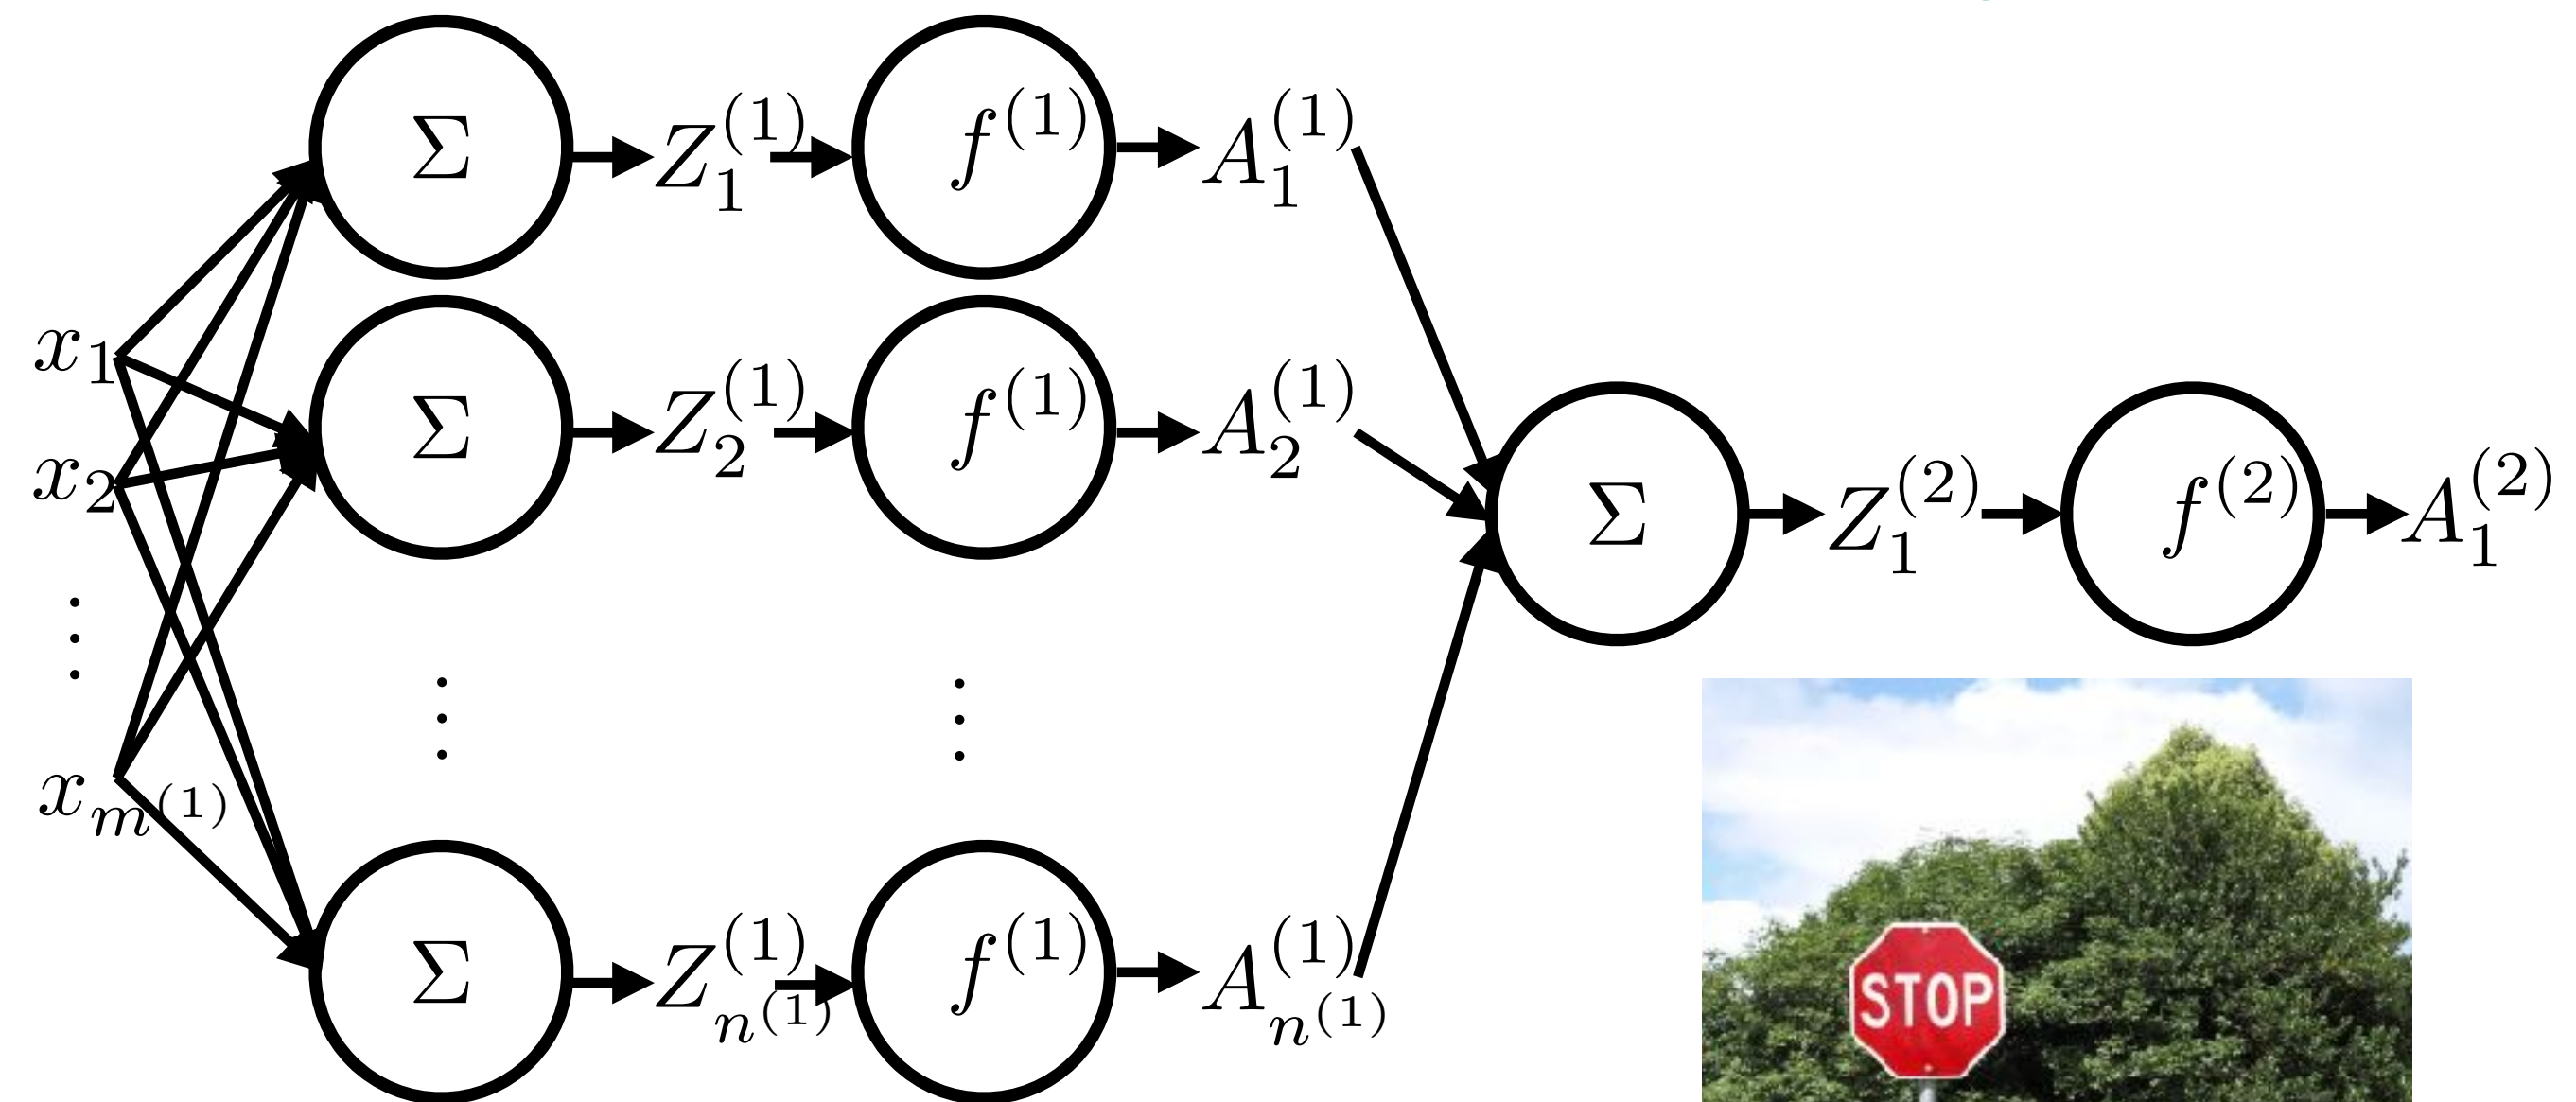
But we know more about images:

- Spatial locality
- Translation invariance



Previous neural nets in this class

- Recall: *Fully connected layer: every input is connected to every output by a weight*




But we know more about images:

- Spatial locality
- Translation invariance



Convolutional Layer: 1D example

A 1D image:




A filter:




$$0 * -1 + 0 * 1 + 1 * -1 = -1$$

After
convolution*:

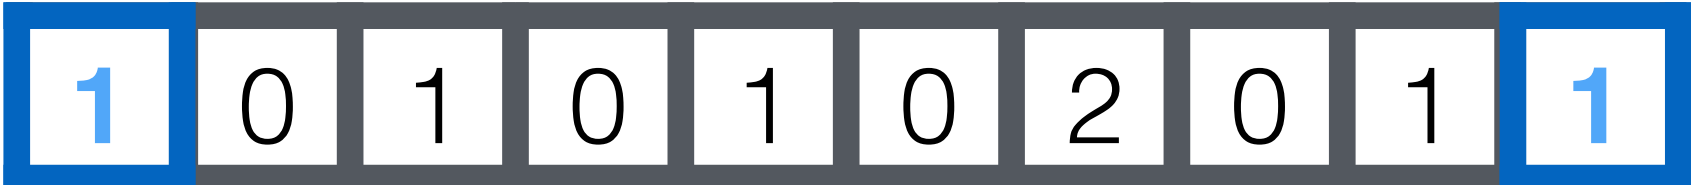


Convolutional Layer: 1D example

A 1D image: 


A filter:  with bias +1

After convolution*: 

After ReLU: 

Convolutional Layer: 1D example

A 1D image: 

A filter:  with bias b

After convolution*: 

After ReLU: 

- How many weights (including bias)? 4
- How many weights (including biases) for fully connected layer with 10 inputs & 10 outputs? $10 \times 11 = 110$

Convolutional Layer: 2D example

A 2D
image:

0	0	0	0	0	0	0
0	1	0	1	0	0	0
0	1	0	1	0	1	0
0	1	1	1	0	0	0
0	1	0	1	0	1	0
0	1	0	1	0	1	0
0	0	0	0	0	0	0

A filter:

-1	-1	-1
-1	1	-1
-1	-1	-1

After
convolution:

0	-4	0	-3	-1
-2	-7	-2	-4	1
-2	-5	-2	-5	-2
-2	-7	-2	-5	0
0	-4	0	-4	0

Convolutional Layer: 2D example

A 2D
image:

0	0	0	0	0	0	0
0	1	0	1	0	0	0
0	1	0	1	0	1	0
0	1	1	1	0	0	0
0	1	0	1	0	1	0
0	1	0	1	0	1	0
0	0	0	0	0	0	0

A filter:

-1	-1	-1
-1	1	-1
-1	-1	-1

After
convolution
& ReLU:

0	0	0	0	0
0	0	0	0	1
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1


A filter:

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9


with bias b

Convolutional Layer: 3D example

A 3D
image:



A filter:




- Tensor: generalization of a matrix
 - E.g. 1D: vector, 2D: matrix

[<https://helpx.adobe.com/photoshop/key-concepts/skew.html>]



Convolutional Layer: multiple filters

An
image:



- Collection of filters in the layer: *filter bank*
- Each resulting image is a *channel*

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- E.g. size 3x3 (“size 3”)

After max pooling:

0	0	1	1
1	1	1	1
1	1	0	0
1	1	0	0

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- E.g. size 3x3 (“size 3”)
- E.g. stride 1

After max pooling:

0	0	1	1
1	1	1	1
1	1	0	0
1	1	0	0

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- E.g. size 3x3 (“size 3”)
- E.g. stride 3

After max pooling:

0	1
1	0

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

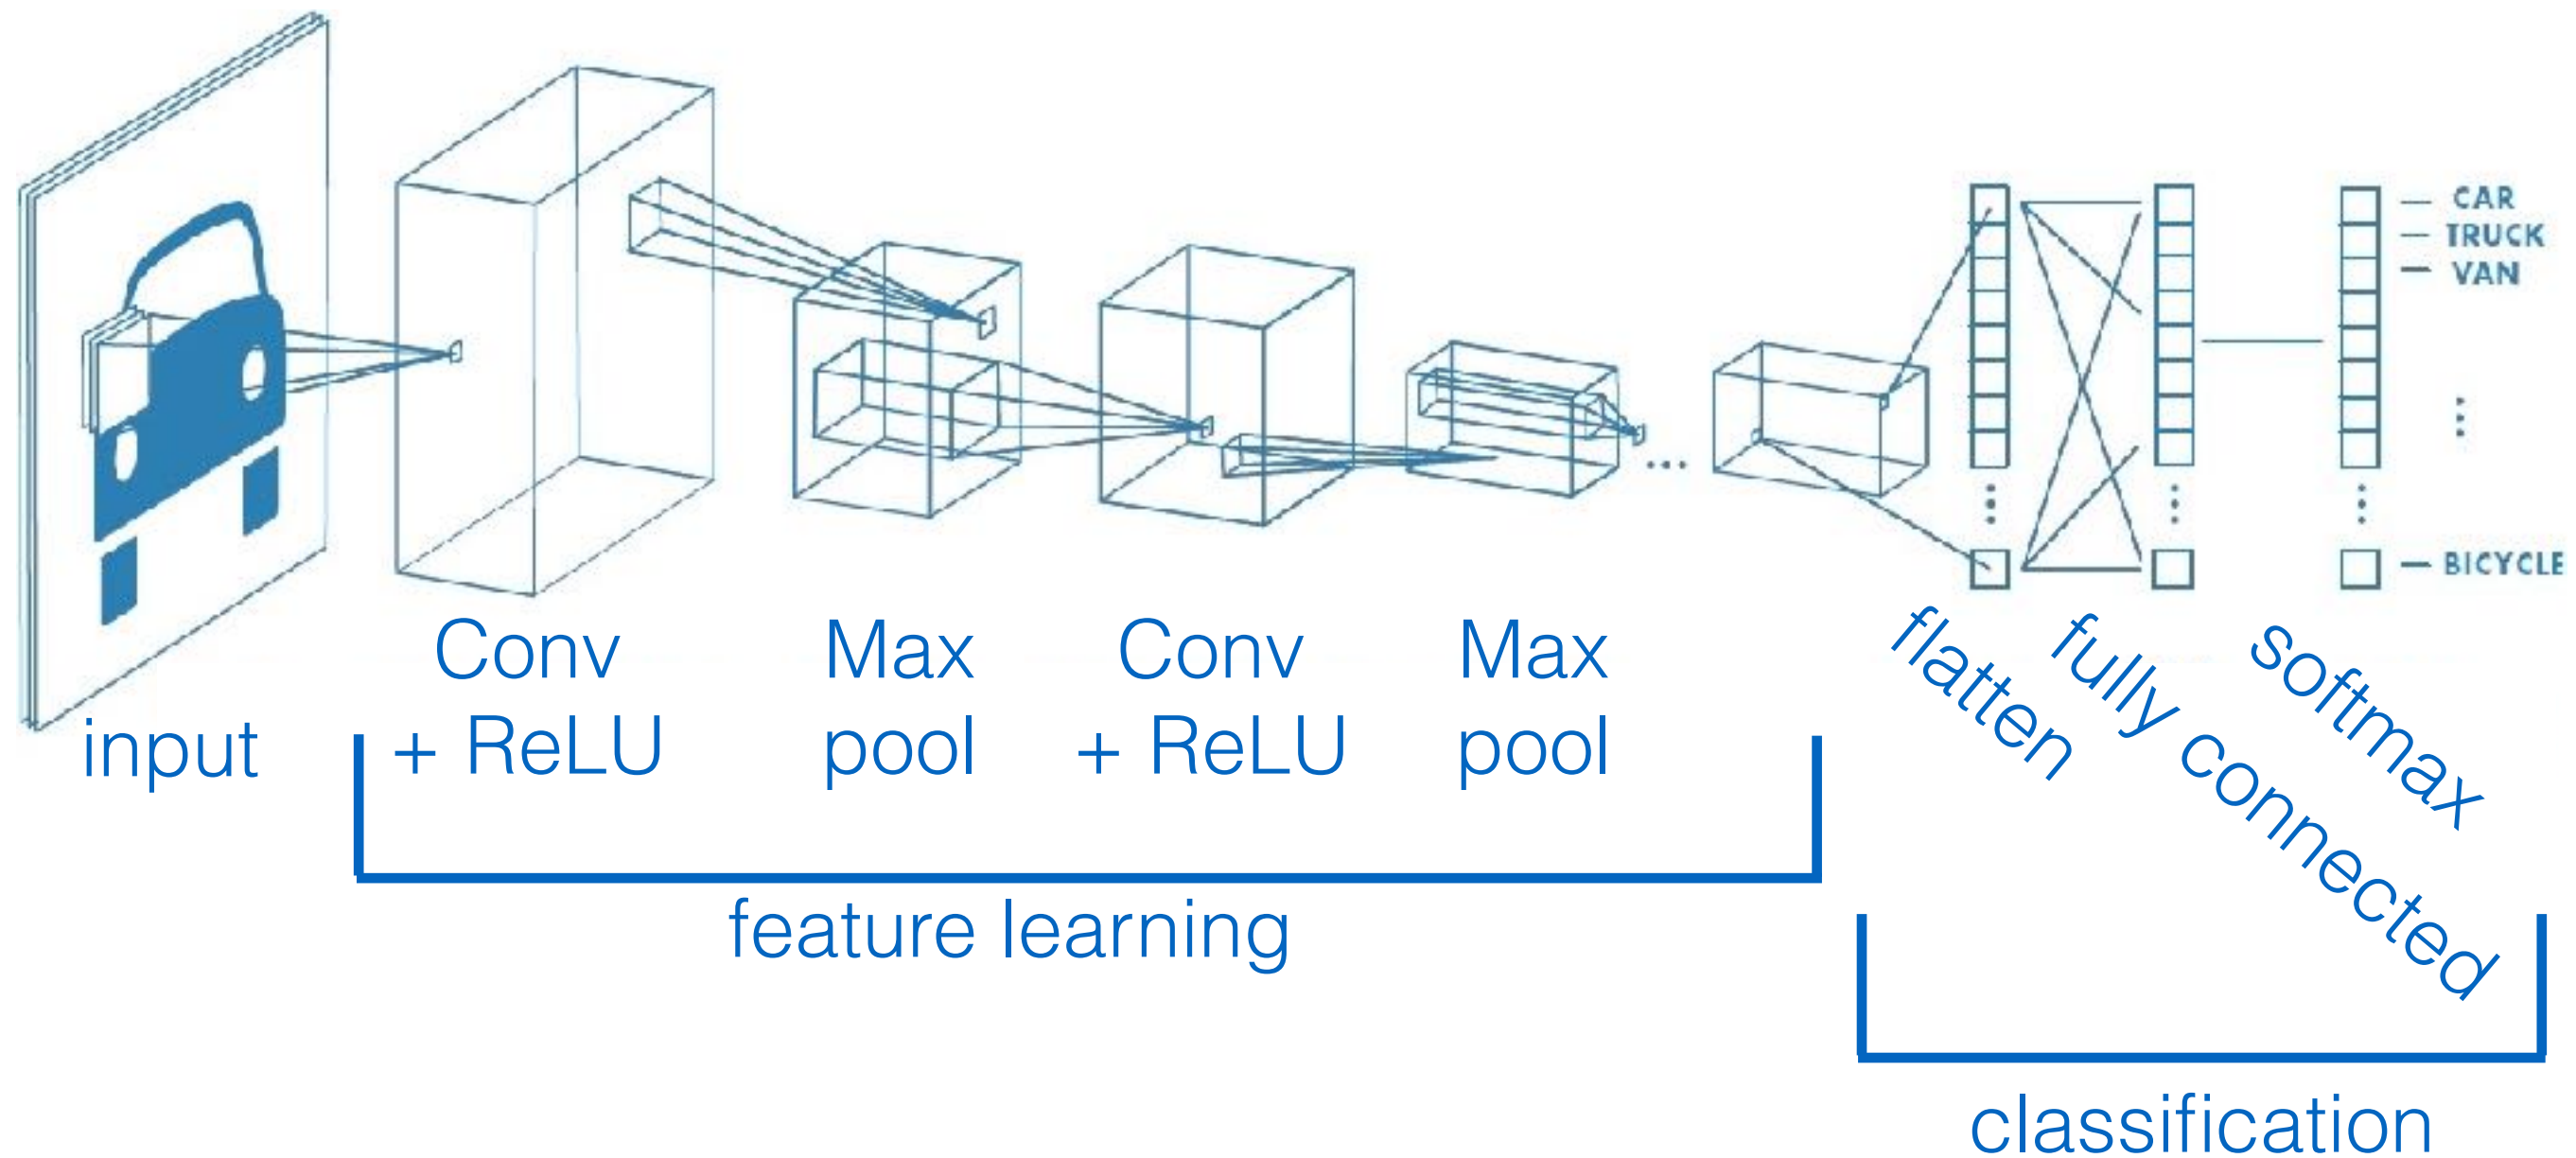
- E.g. size 3x3 (“size 3”)
- E.g. stride 3

After max pooling:

0	1
1	0

- Can use stride with filters too
- No weights in max pooling

CNNs: typical architecture




x

$\text{NN}(x; W, W_0)$

A familiar pattern

1. Choose how to predict label (given features & parameters)
2. Choose a loss (between guessed label & actual label)
3. Choose parameters by trying to minimize the training loss



CNNs: a taste of backpropagation

Regression. 1 filter: size 3 & padding; $x^{(j)}$ dimension: 5x1

- Forward pass:

$$Z_i^1 = (W^1)^\top X_{[i-1,i,i+1]} \quad Z^1: 5 \times 1$$

$$A_i^1 = \text{ReLU}(Z_i^1) \quad A^1: 5 \times 1$$

$$A^2 = (W^2)^\top A^1 \quad A^2: 1 \times 1$$

$$L(A^2, y) = (A^2 - y)^2 \quad \text{Loss: } 1 \times 1$$

- Part of the derivative for SGD:

$$\frac{\partial \text{loss}}{\partial W^1} = \frac{\partial Z^1}{\partial W^1} \cdot \frac{\partial A^1}{\partial Z^1} \cdot \frac{\partial \text{loss}}{\partial A^1}$$

$3 \times 1 \quad \quad 3 \times 5 \quad \quad 5 \times 5 \quad \quad 5 \times 1$

$$Z_2^1 = W_1^1 X_1 + W_2^1 X_2 + W_3^1 X_3$$

$$\frac{\partial Z^1}{\partial W^1} = \begin{bmatrix} \overset{Z_1^1}{\boxed{X_0}} & \overset{Z_2^1}{X_1} & \overset{Z_3^1}{X_2} & \overset{Z_4^1}{X_3} & \overset{Z_5^1}{X_4} \\ X_1 & X_2 & X_3 & X_4 & X_5 \\ X_2 & X_3 & X_4 & X_5 & \boxed{X_6} \end{bmatrix} \begin{matrix} W_1^1 \\ W_2^1 \\ W_3^1 \end{matrix}$$